




Sistemas Operativos

Threads

- These slides and notes are based on the contents of the books:
 - Abraham Silberschatz, "Operating System Concepts", 10th Edition, Wiley, 2018;
 - William Stallings, "Operating Systems: Internals and Design Principles", 9th Edition, Pearson, 2017;
 - Andrew S. Tanenbaum, "Modern Operating Systems", 4th Edition, Pearson Education, 2014;
- The respective copyrights belong to their owners.



- **Motivation**

- **Thread Concept**

- Examples
- Benefits

- **Thread Fields**

- Program counter
- Registers
- Stack
- State

- **Multithreading Models**

- Many-to-one Model (M-to-1)
- One-to-one Model (1-to-1)
- Many-to-many Model (M-to-N)
- Two-level Model

- **Threading Libraries**

- Pthreads, Win32, Java

- **Pthreads**

- Management

- `int pthread_create()`
- `void pthread_exit()`
- `int pthread_join()`
- `int pthread_cancel()`

- Compilation

- `-D_REENTRANT`
- `-lpthread`

- **References**


- **Motivation**

- Process switching is an extremely heavy operation
 - necessary to remap address spaces
 - necessary to manage all information associated to processes
- **Thread**
 - It's a flow of execution inside a program
 - The address space is the same (changing a global variable in a thread affects all others)
 - No context switching between processes
 - Commutation among threads is **very fast**
 - Communication among threads is easy to do and fast

- **Motivation**
- **Thread Concept**
 - Examples
 - Benefits
- **Thread Fields**
 - Program counter
 - Registers
 - Stack
 - State
- **Multithreading Models**
 - Many-to-one Model (M-to-1)
 - One-to-one Model (1-to-1)
 - Many-to-many Model (M-to-N)
 - Two-level Model
- **Threading Libraries**
 - Pthreads, Win32, Java
- **Pthreads**
 - Management
 - `int pthread_create()`
 - `void pthread_exit()`
 - `int pthread_join()`
 - `int pthread_cancel()`
 - Compilation
 - `-D_REENTRANT`
 - `-lpthread`
- **References**

• Thread Concept

- One way of looking at a **process** is that it is a way to group related resources together.
- A process has an address space containing:
 - program text and data
 - other resources such as open files, child processes, pending alarms, signal handlers, accounting information, and more



process stack - which contains temporary data (such as function parameters, return addresses, and local variables)

data section - which contains global variables

heap - which is memory that is dynamically allocated during process run time (e.g. as a result of *malloc()*).

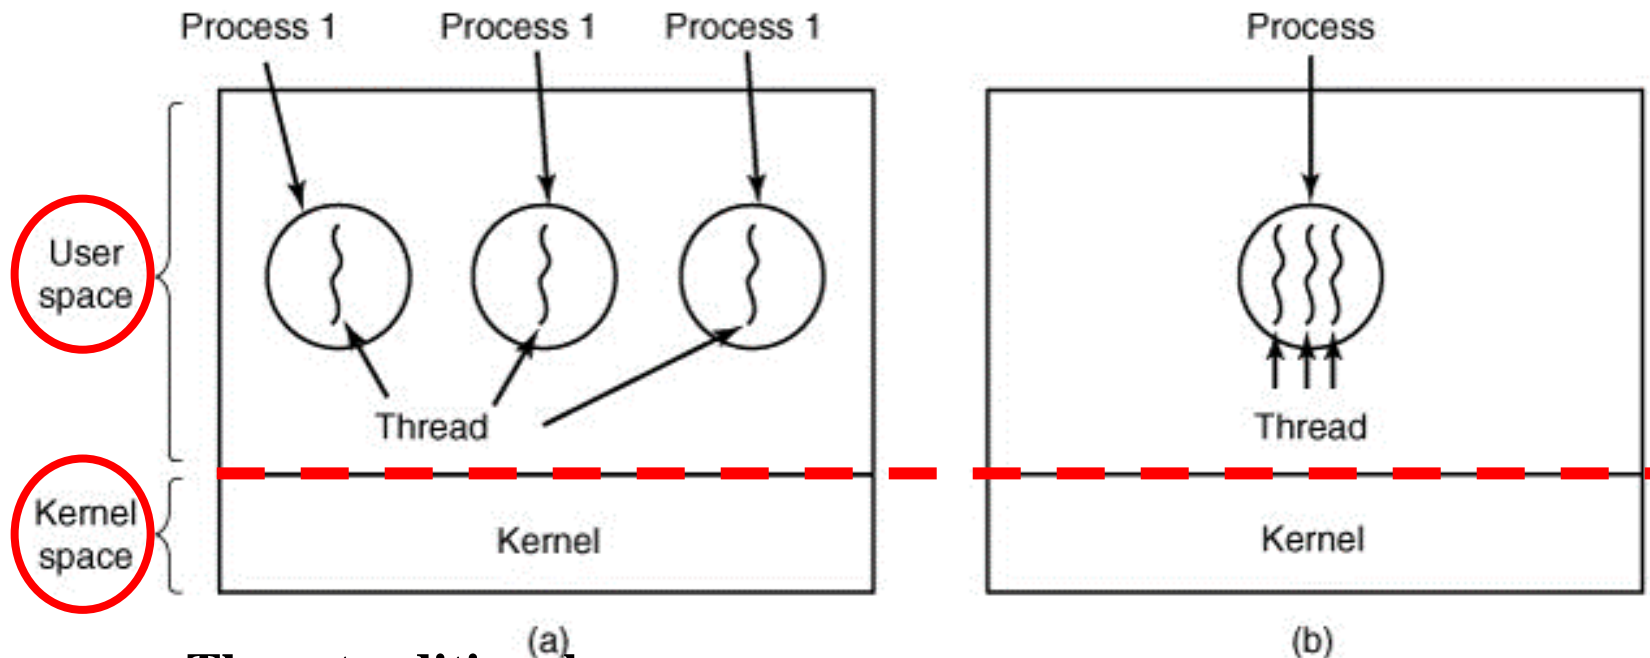
text - where the code of the program goes

- **Thread Concept**

- The other concept a **process** has is a thread of execution which contains a:
 - a program counter (that keeps track of which instruction to execute next)
 - registers (which hold its current working variables).
 - a stack (which contains the execution history with one frame for each procedure called but not yet returned from).
- Although a thread must execute in some process, the thread and its process are different concepts and can be treated separately.
- **Processes are used to group resources together; threads are the entities scheduled for execution on the CPU.**

• Thread Concept

- threads to allow multiple executions to take place in the same process environment, to a large degree independent of one another.




Three traditional processes.
Each process has its own address space and a single thread of control. Each of the threads operates in a different address space.

Single process with three threads of control.
All three threads share the same address space

• Thread Concept

} = instruction trace



William Stallings, "Operating Systems: Internals and Design Principles"

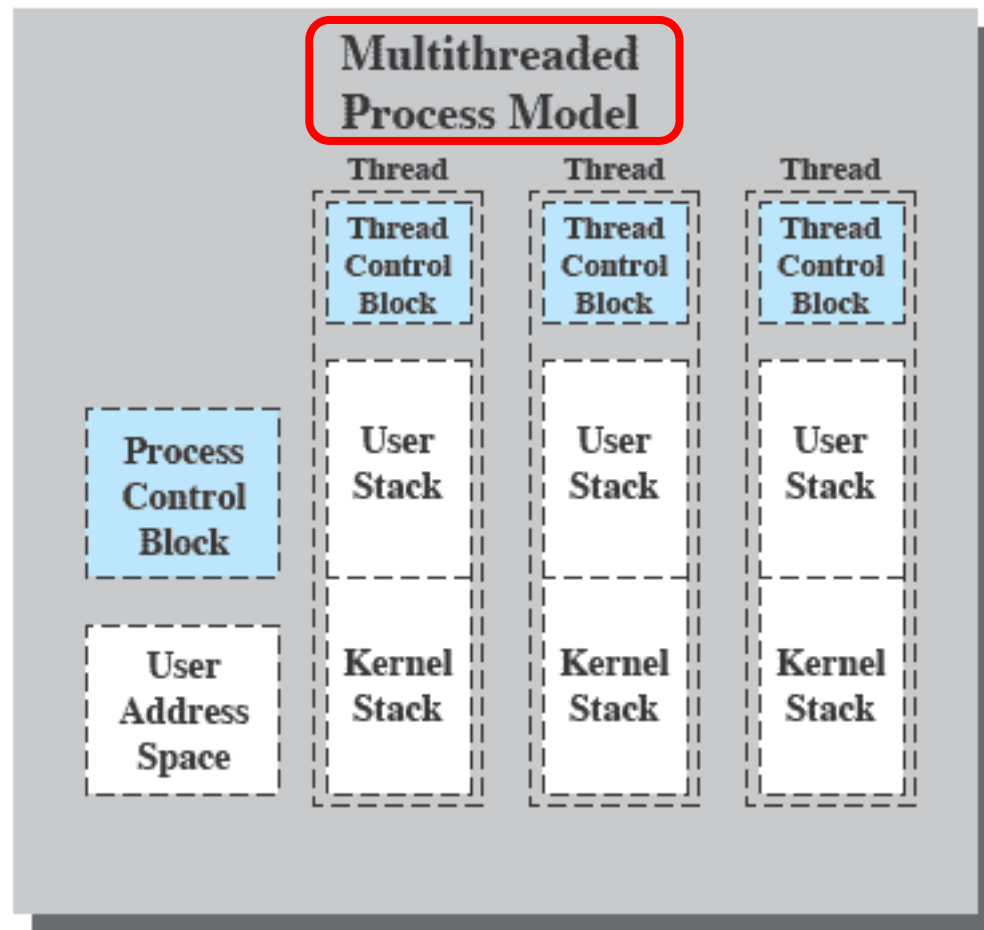
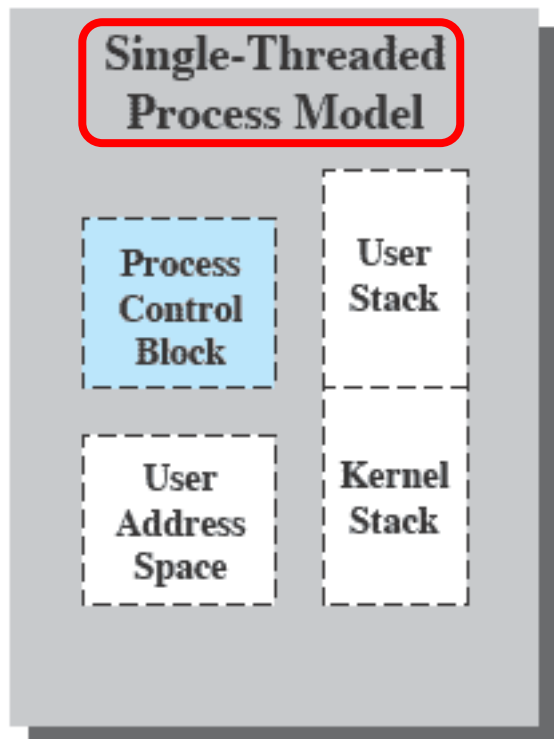
- Thread Concept

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional <u>UNIX</u> implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	<u>Windows NT, Solaris, Linux, OS/2, OS/390, MACH</u>

William Stallings, “Operating Systems: Internals and Design Principles”


- **Thread Concept**

One way to view a **thread** is as an independent program counter operating within a process.

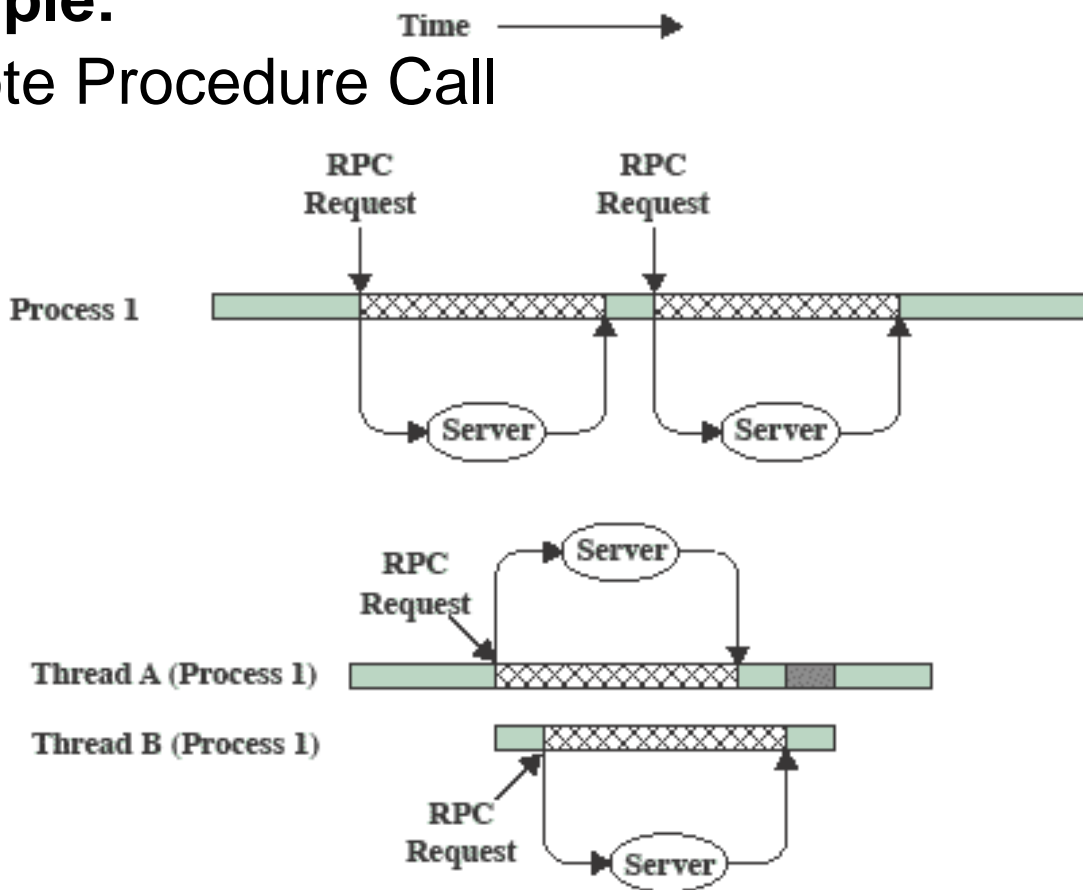


- **Example:**
- Web Browser:
 - web pages may contain multiple small images.
 - For each image on a web page, the browser must set up a separate connection to the page's home site and request the image. A great deal of time is spent establishing and releasing all these connections.
 - By having multiple threads within the browser, many images can be requested at the same time, greatly speeding up performance in most cases since with small images, the set-up time is the limiting factor, not the speed of the transmission line.

- **Example:**
- **Web Server:**






- **Example:**
- Remote Procedure Call

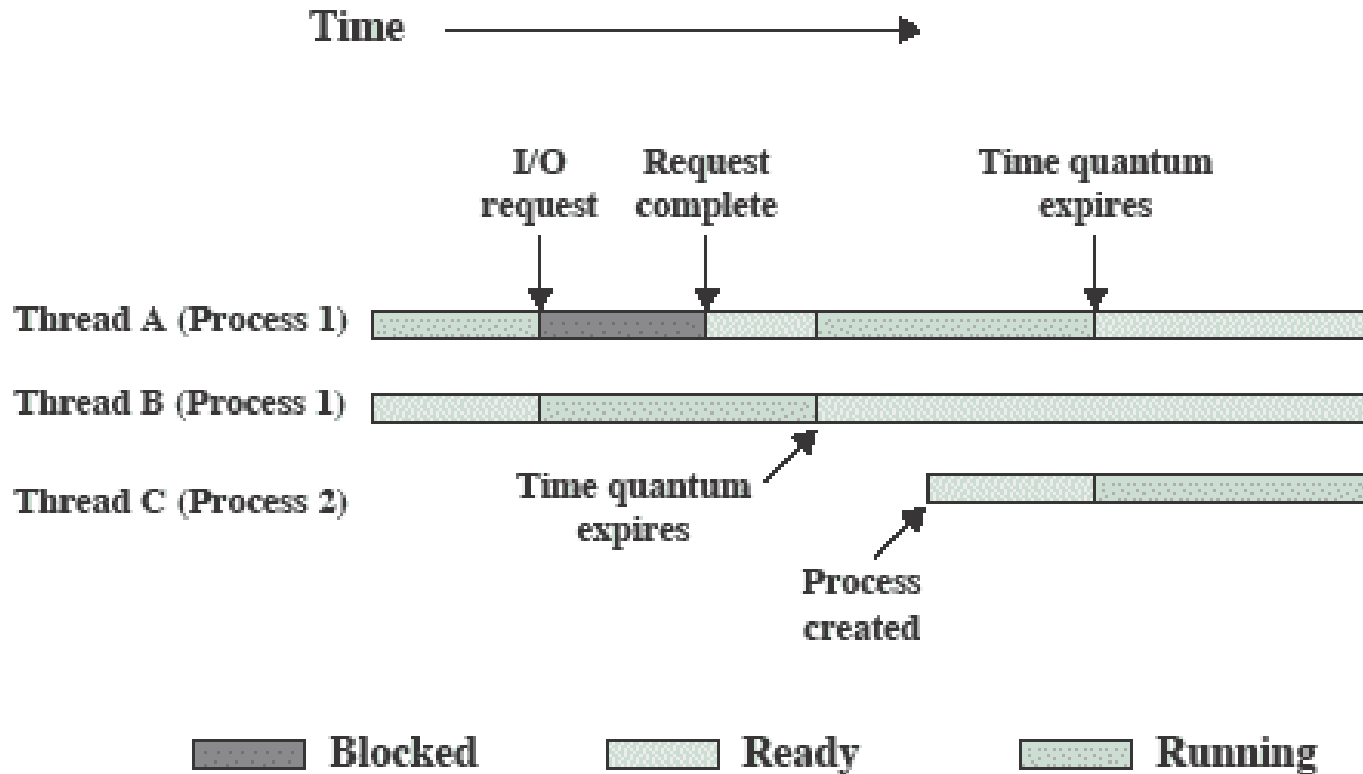


**RPC using
single thread**

**RPC using one
thread per
server (on a
uniprocessor)**

-  Blocked, waiting for response to RPC
-  Blocked, waiting for processor, which is in use by Thread B
-  Running

• Example:



Multithreading on a uniprocessor

- **Benefits**

- **very light weight compared to processes:**
 - Light context switches
 - Fast to create and terminate
 - Fast to synchronize
 - Share resources
- **Responsiveness:** multithreading may allow a program to continue running even if part of it is blocked or is performing a lengthy operation.
- **Resource sharing/economy:** threads share the memory and the resources of the process to which they belong.
- **Utilization of multiprocessor architectures:** threads may be running in parallel on different processors.

- **Motivation**
- **Thread Concept**
 - Examples
 - Benefits
- **Thread Fields**
 - Program counter
 - Registers
 - Stack
 - State
- **Multithreading Models**
 - Many-to-one Model (M-to-1)
 - One-to-one Model (1-to-1)
 - Many-to-many Model (M-to-N)
 - Two-level Model
- **Threading Libraries**
 - Pthreads, Win32, Java
- **Pthreads**
 - Management
 - `int pthread_create()`
 - `void pthread_exit()`
 - `int pthread_join()`
 - `int pthread_cancel()`
 - Compilation
 - `-D_REENTRANT`
 - `-lpthread`
- **References**

- **Thread Fields**


- When multiple threads are present in the same address space, some fields are not per process, but per thread, so a separate thread table is needed, with one entry per thread.
- Among the **per-thread items**:
 - program counter needed because threads, like processes, can be suspended and resumed
 - registers needed because when threads are suspended, their registers must be saved
 - stack to store local variables
 - state because threads, like processes, can be in running, ready, or blocked state.

• Thread Fields

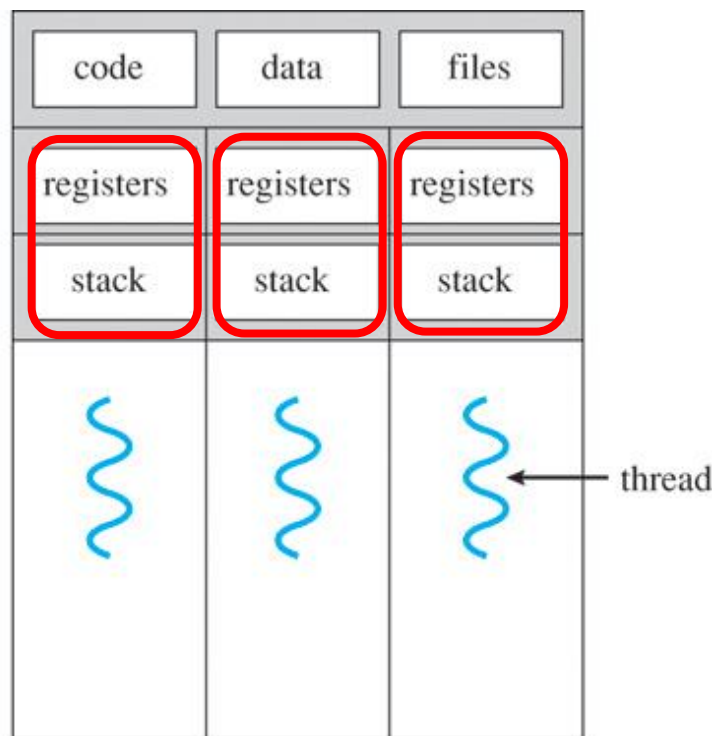
Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

• Thread Fields

- **Threads share the same address space and resources!**
Changing one thing in one thread affects all others!
- It is the responsibility of the programmer to assure the correctness in the concurrent access to data and resources!

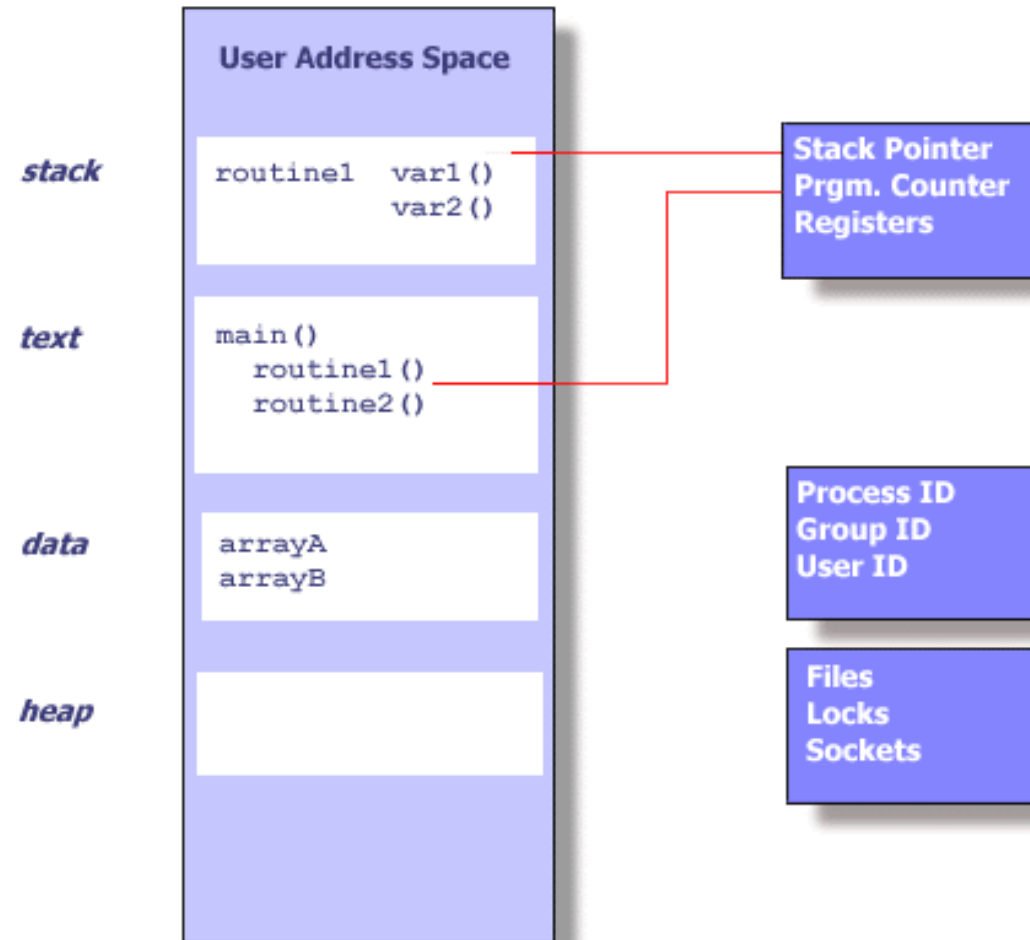


single-threaded process

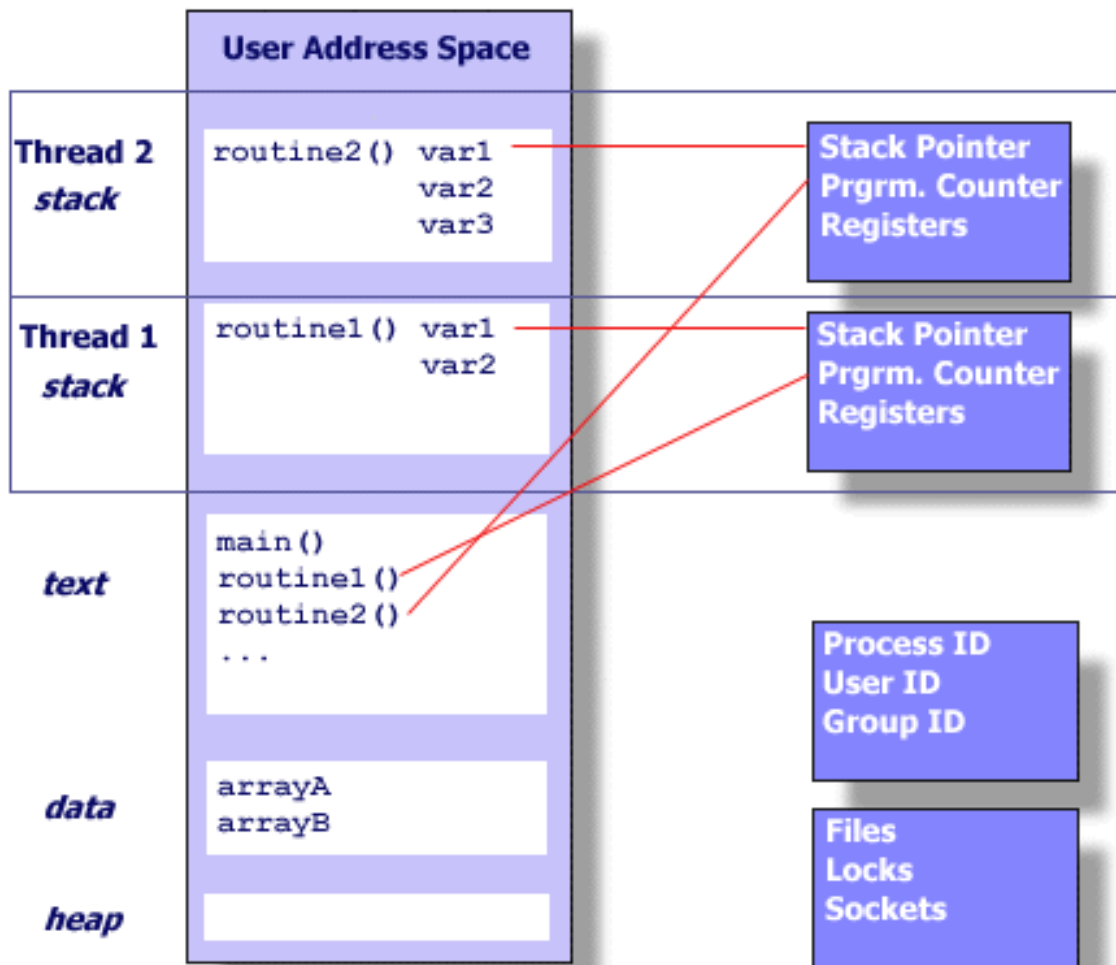


multithreaded process

• Thread Fields



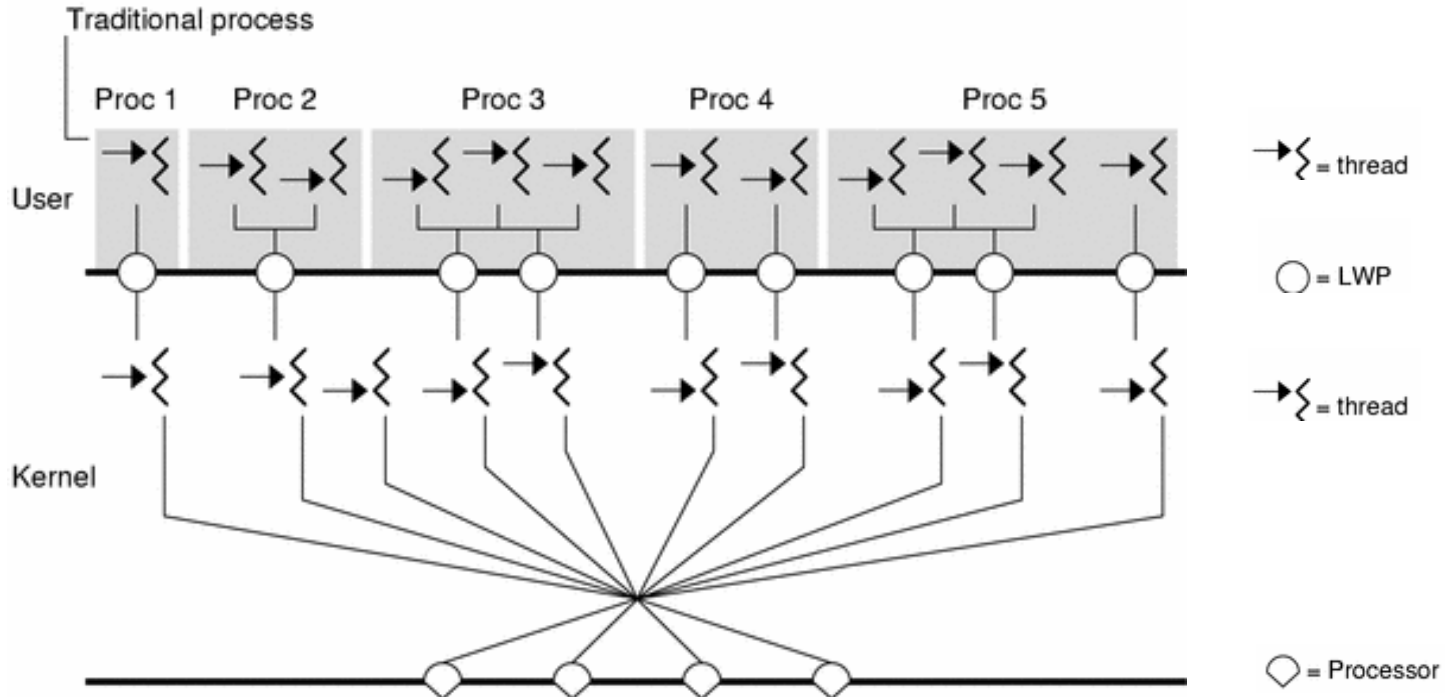
• Thread Fields



- **Motivation**
- **Thread Concept**
 - Examples
 - Benefits
- **Thread Fields**
 - Program counter
 - Registers
 - Stack
 - State
- **Multithreading Models**
 - Many-to-one Model (M-to-1)
 - One-to-one Model (1-to-1)
 - Many-to-many Model (M-to-N)
 - Two-level Model
- **Threading Libraries**
 - Pthreads, Win32, Java
- **Pthreads**
 - Management
 - `int pthread_create()`
 - `void pthread_exit()`
 - `int pthread_join()`
 - `int pthread_cancel()`
 - Compilation
 - `-D_REENTRANT`
 - `-lpthread`
- **References**

• Multithreading Models

- Support for threads may be provided either at the user level, for user threads, or by the kernel, for kernel threads.
 - **User threads** are supported above the kernel and are managed without kernel support
 - **Kernel threads** are supported and managed directly by the operating system.




- **Multithreading Models**

- There must exist a relationship between user threads and kernel threads
 - **Many-to-one Model (M-to-1):**
 - maps many user-level threads to one kernel thread.
 - **One-to-one Model (1-to-1):**
 - maps each user thread to a kernel thread.
 - **Many-to-many Model (M-to-N):**
 - multiplexes many user-level threads to a smaller or equal number of kernel threads.
 - **Two-level Model:**
 - multiplexes many user-level threads to a smaller or equal number of kernel threads but also allows a user-level thread to be bound to a kernel thread.

• Multithreading Models

- Many-to-one Model (M-to-1 model)


- Thread management is done by the thread library in user space, so it is efficient
- The kernel is completely unaware of the existence of threads
- The entire process will block if a thread makes a blocking system call.
- Because only one thread can access the kernel at a time, multiple threads are unable to run in parallel on multiprocessors.



• Multithreading Models

- One-to-one Model (1-to-1 model)


- Threads are implemented exclusively in kernel space
- The kernel does all the scheduling of threads
- It provides more concurrency than the many-to-one model by allowing another thread to run when a thread makes a blocking system call
- Allows multiple threads to run in parallel on multiprocessors.
- The only drawback to this model is that creating a user thread requires creating the corresponding kernel thread.



• Multithreading Models

- Many-to-many Model (M-to-N model)


- A number of kernel threads can map to a different number of user threads.
- The number of kernel threads may be specific to either a particular application or a particular machine (an application may be allocated more kernel threads on a multiprocessor than on a uniprocessor).
- When a thread performs a blocking system call, the kernel can schedule another thread for execution.







- **Multithreading Models**

- **Two level model**

- multiplexes many user-level threads to a smaller or equal number of kernel threads, but also allows a user-level thread to be bound to a kernel thread



• Multithreading Models



- **Motivation**
- **Thread Concept**
 - Examples
 - Benefits
- **Thread Fields**
 - Program counter
 - Registers
 - Stack
 - State
- **Multithreading Models**
 - Many-to-one Model (M-to-1)
 - One-to-one Model (1-to-1)
 - Many-to-many Model (M-to-N)
 - Two-level Model
- **Threading Libraries**
 - Pthreads, Win32, Java
- **Pthreads**
 - Management
 - `int pthread_create()`
 - `void pthread_exit()`
 - `int pthread_join()`
 - `int pthread_cancel()`
 - Compilation
 - `-D_REENTRANT`
 - `-lpthread`
- **References**

• Threading Libraries

- A thread library provides the programmer an API for creating and managing threads. There are two primary ways of implementing a thread library:
 - to provide a library entirely in user space with no kernel support. All code and data structures for the library exist in user space. This means that invoking a function in the library results in a local function call in user space and not a system call.
 - to implement a kernel-level library supported directly by the operating system. In this case, code and data structures for the library exist in kernel space. Invoking a function in the API for the library typically results in a system call to the kernel.

- **Threading Libraries**

- Three main thread libraries are in use:

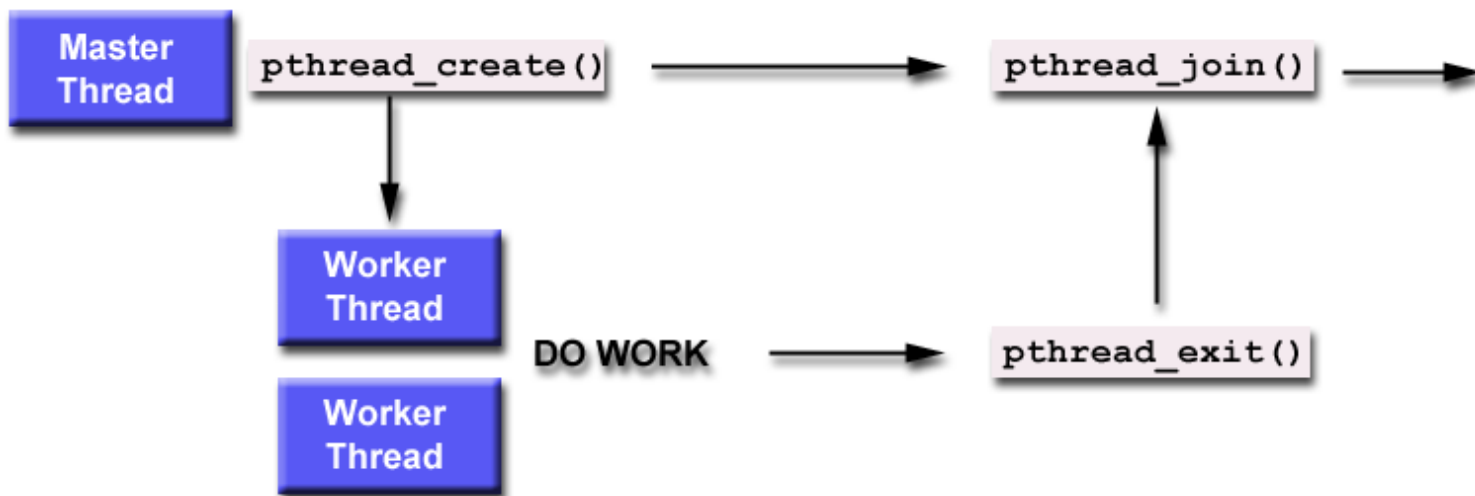
- **POSIX Threads (Pthreads)**: the threads extension of the POSIX standard, may be provided as either a user- or kernel-level library.
 - **Win32 Threads**: kernel-level library available on Windows systems.
 - **Java Threads**: because in most instances the JVM is running on top of a host operating system, the Java thread API is typically implemented using a thread library available on the host system.
 - **Windows systems**:
 - typically implemented using the Win32 API;
 - **UNIX and Linux systems**:
 - often use Pthreads

- **Motivation**
- **Thread Concept**
 - Examples
 - Benefits
- **Thread Fields**
 - Program counter
 - Registers
 - Stack
 - State
- **Multithreading Models**
 - Many-to-one Model (M-to-1)
 - One-to-one Model (1-to-1)
 - Many-to-many Model (M-to-N)
 - Two-level Model
- **Threading Libraries**
 - Pthreads, Win32, Java
- **Pthreads**
 - Management
 - `int pthread_create()`
 - `void pthread_exit()`
 - `int pthread_join()`
 - `int pthread_cancel()`
 - Compilation
 - `-D_REENTRANT`
 - `-lpthread`
- **References**

- **Pthreads**
- **Management**

- Pthreads refers to the **POSIX standard (IEEE 1003.1c)** defining an API for thread creation and synchronization:

- *int pthread_create()*
- *void pthread_exit()*
- *int pthread_join()*
- *int pthread_cancel()*



- Pthreads
 - Example 1

```
/Temp/so$ ./thread_demo
Hello, I'm a Slooowww THREAD!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!
Hello, I'm a Slooowww THREAD!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!
Hello, I'm a Slooowww THREAD!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!
Hello, I'm a Slooowww THREAD!
Hello, I'm a FAST Thread!
Hello, I'm a FAST Thread!

/Temp/so$ _
```

```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>
```

```
void* slow_thread() {
    while (1) {
        printf("Hello, I'm a Slooowww THREAD!\n");
        sleep(3);
    }
}
```

```
void* fast_thread() {
    while (1) {
        printf("Hello, I'm a FAST Thread!\n");
        sleep(1);
    }
}
```

```
int main() {
    pthread_t thr_slow, thr_fast;

    pthread_create(&thr_slow, NULL, slow_thread, NULL);
    pthread_create(&thr_fast, NULL, fast_thread, NULL);

    pthread_exit(NULL);
    return 0;
}
```

summation of a nonnegative integer in a separate thread.

• Pthreads

- Example 2

Separate threads begin execution in a specified function which is in this case *runner()*

```
#include <pthread.h>
#include <stdio.h>
```

Pthreads programs must include the pthread.h header file.

```
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */
```

```
int main(int argc, char *argv[])
```

The statement `pthread_t tid` declares the identifier for the thread we will create.

```
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */
```

The `pthread_attr_t attr` declaration represents the attributes for the thread. Namely, the stack size and scheduling information.

```
    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

• Pthreads

- Example 2

At this point, the program has two threads: the initial (or parent) thread in *main ()* and the summation (or child) thread performing the summation operation in the *runner ()* function.

the parent thread will wait for it to complete by calling the *pthread_join ()* function

Both threads share the global data *sum*.

The summation thread will complete when it calls the function *pthread_exit ()*

```
/* get the default attributes */  
pthread_attr_init(&attr);  
/* create the thread */  
pthread_create(&tid, &attr, runner, argv[1]);  
/* wait for the thread to exit */  
pthread_join(tid, NULL);
```

```
printf("sum = %d\n", sum);  
}
```

```
/* The thread will begin control in this function */  
void *runner(void *param)  
{  
    int i, upper = atoi(param);  
    sum = 0;  
  
    for (i = 1; i <= upper; i++)  
        sum += i;  
  
    pthread_exit(0);  
}
```

set the attributes in the function call *pthread_attr_init (&attr)*. Because we did not explicitly set any attributes, we use the default attributes provided.

A separate thread is created with the *pthread_create ()* function call, passing the thread identifier, the attributes for the thread, the name of the function where the new thread will begin execution (in this case, the *runner ()* function) and the integer parameter that was provided on the command line, *argv [1]*

- Pthreads
 - Example 3

```
// Ids used by the threads
pthread_t my_thread[N];
int id[N];
```

```
// worker thread
void* worker(void* idp) {
    int my_id = *((int*) idp);

    printf("Hello, I'm thread %d\n", my_id);
    sleep(rand()%3);
    printf("Hello, I'm thread %d, going away!\n", my_id);

    pthread_exit(NULL);
    return NULL;
}
```

```
int main()
{
    // Creates N threads
    for (int i=0; i<N; i++) {
        id[i] = i;
        pthread_create(&my_thread[i], NULL, worker, &id[i]);
    }

    // waits for them to die
    for (int i=0; i<N; i++)
        pthread_join(my_thread[i], NULL);

    return 0;
}
```

- **Pthreads**
- **Compilation**
 - use these library calls we must include the file `pthread.h` and link with the threads library using `-lpthread`
 - Re-entrant code can be called more than once, whether by different threads or by nested invocations in some way, and still work correctly (e.g. `errno` variable). The macro `REENTRANT` must be defined in the compilation
 - Example:
 - `$ gcc -D_REENTRANT thread1.c -o thread1 -lpthread`

- Abraham Silberschatz, "Operating System Concepts", 10th Edition, Wiley, 2018
- William Stallings, "Operating Systems: Internals and Design Principles", 9th Edition, Pearson, 2017
- Andrew S. Tanenbaum, "Modern Operating Systems", 4th Edition, Pearson Education, 2014