




Sistemas Operativos

Processes

- These slides and notes are based on the contents of the books:
 - Abraham Silberschatz, "Operating System Concepts", 10th Edition, Wiley, 2018;
 - William Stallings, "Operating Systems: Internals and Design Principles", 9th Edition, Pearson, 2017;
 - Andrew S. Tanenbaum, "Modern Operating Systems", 4th Edition, Pearson Education, 2014;
- The respective copyrights belong to their owners.



- **Process Concept**

- process stack; data section; heap; text

- **Process State**

- New; Running; Waiting; Ready; Terminated

- **Process Control Block (PCB)**

- **Scheduling Queues**

- Job Queue
- Ready Queue
- Device Queue

- **Schedulers**

- Long-term scheduler
- Short-term scheduler

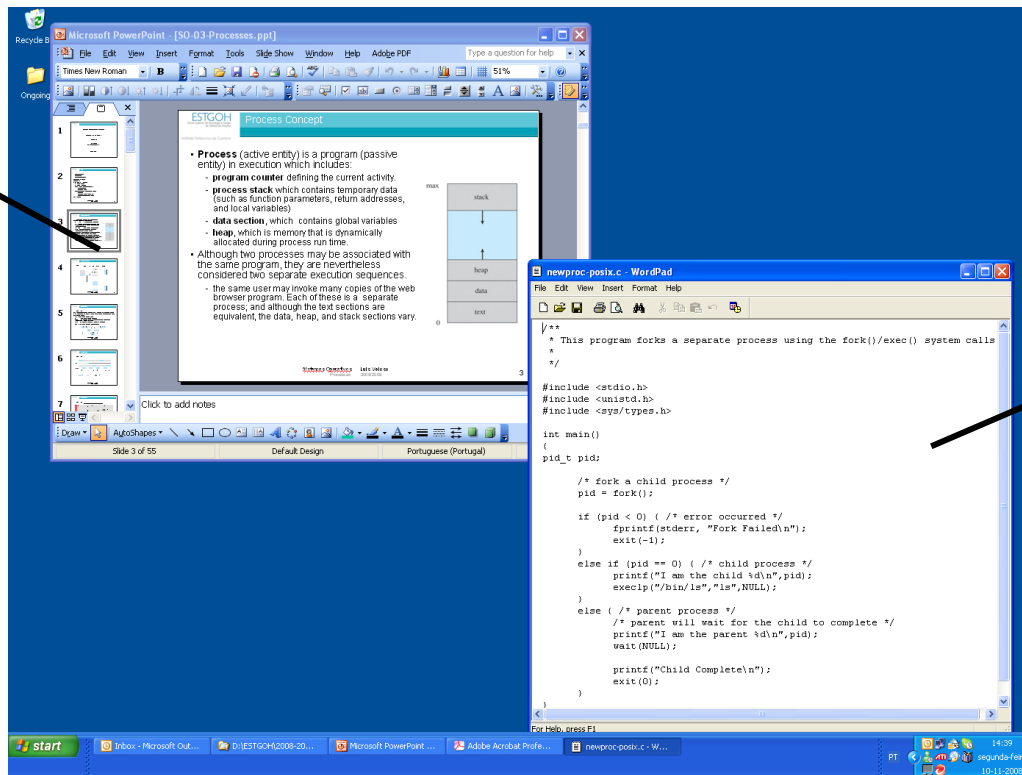
- **Operations in Processes**

- Process Creation
 - fork(), exec(), wait()
- Process Termination
 - exit(), wait()

- **References**

- **Process** (active entity) is a program (passive entity) in execution

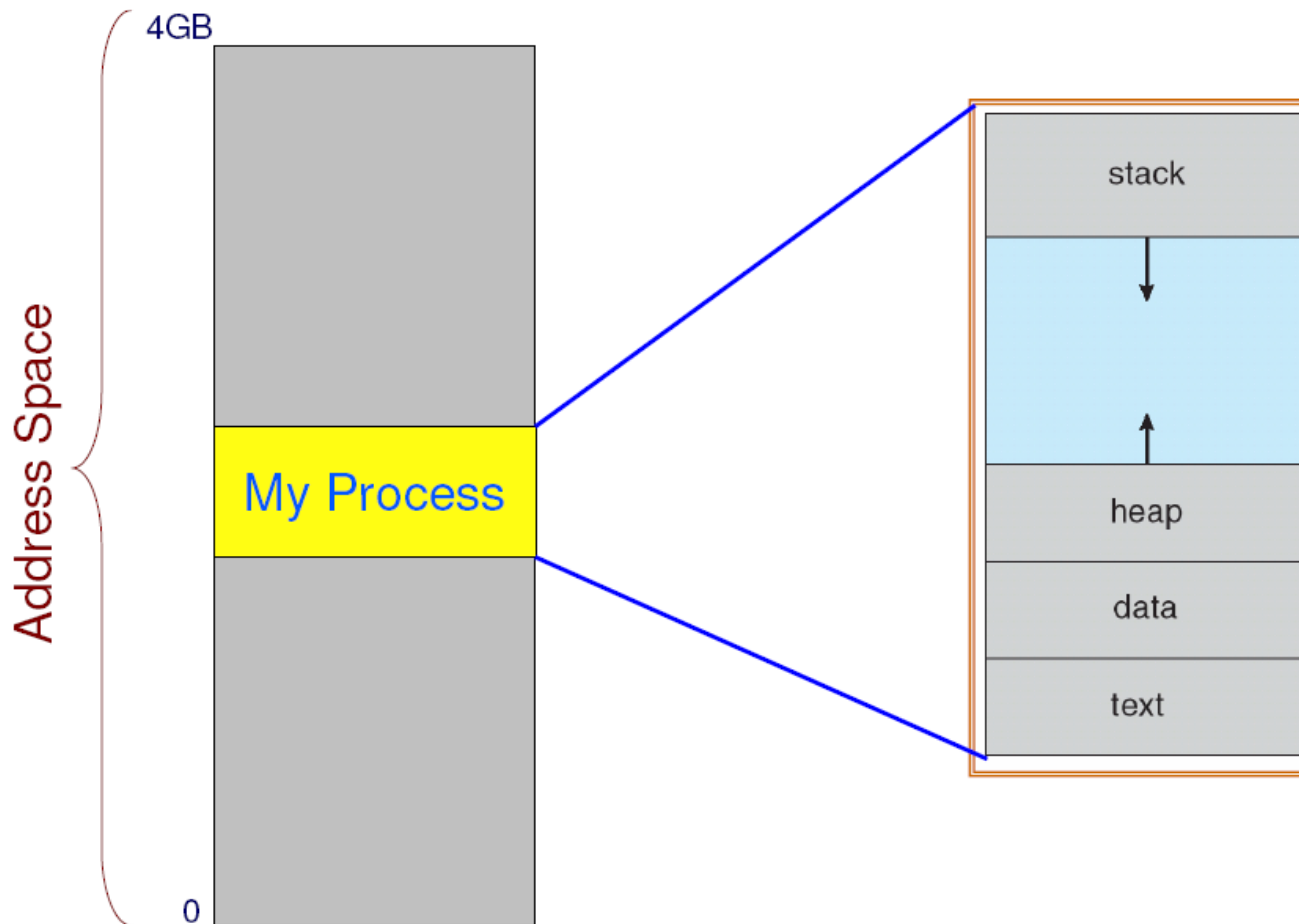
Powerpoint
(1 process)




Wordpad
(1 process)

- Program = Binary Image (Executable)
- Process = The living “Image” of a program running

- **Process** (active entity) is a program (passive entity) in execution



- **Process** is a program in execution which includes:
 - **process stack:** contains temporary data (e.g. as function parameters, return addresses, and local variables)
 - **data section:** contains global variables
 - **heap:** memory that is dynamically allocated during process run time (e.g. as a result of *malloc()*).
 - **text:** where the code of the program goes
- Although two processes may be associated with the same program, they are nevertheless considered two separate execution sequences.



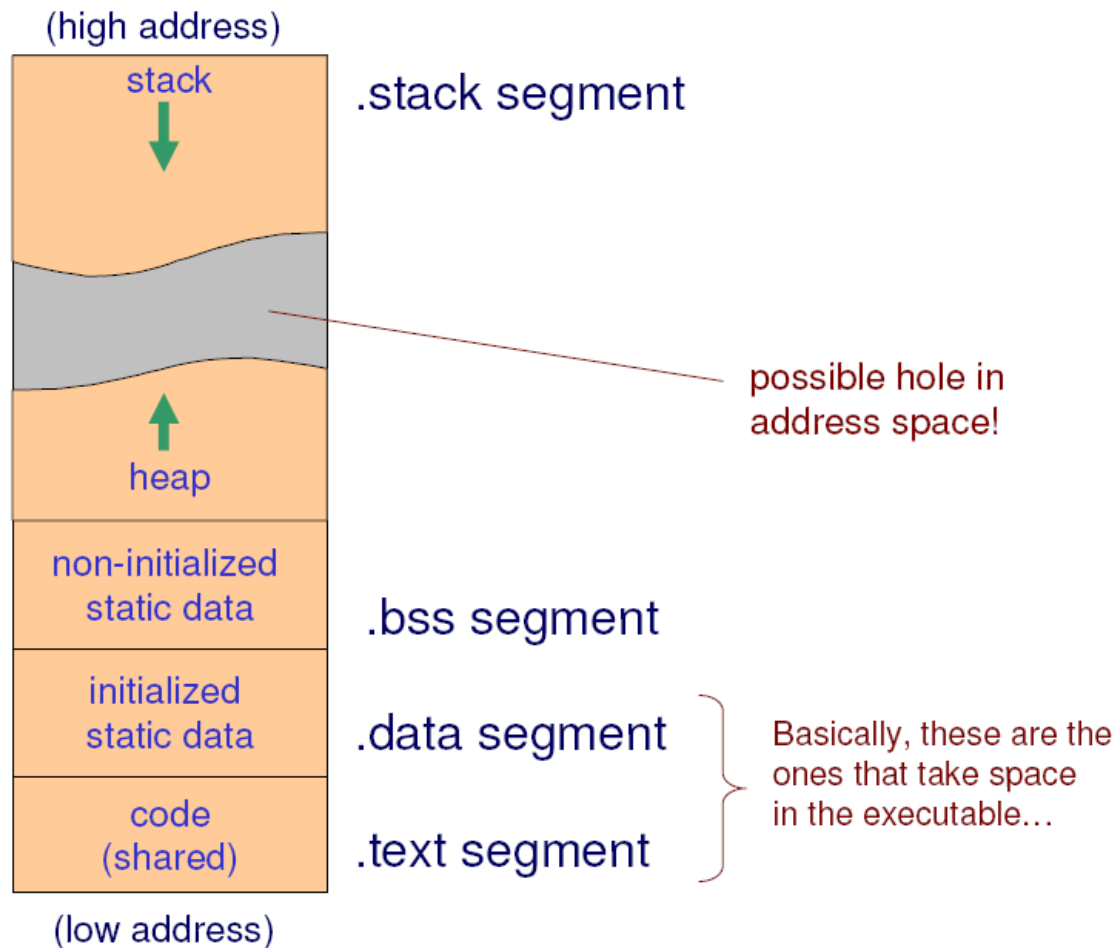
Abraham Silberschatz, "Operating System Concepts"

- **Process**

- **“Process Memory Image” in Unix:**

- **Text:** where the code of the program goes
 - called **.text** section
 - **Data:** where global and static variables are:
 - **.data** section:
 - global and static initialized variables to non-zero
 - **.bss** section:
 - global and static non-initialized variables (or initialized to 0)
 - **Heap:** where all dynamically allocated memory is set (e.g. as a result of malloc())
 - **Stack:** where all automatic variables exist:
 - variables that are automatically created and destroyed in methods
 - it grows down

- **Process**
 - **“Process Memory Image” in Unix:**



• Process

- “Process Memory Image” in Unix:

```
#define KB (1024)
#define MB (1024*1024)

char buf[10*MB];
char command[KB] = "command?";
int n_lines = 0;
int n_tries = 20;
int total;

int f(int n) {
    int result;
    static int number_calls = 0;

    ++number_calls; result = n*n;
    return result;
}

int main() {
    int x = 5;

    printf("f(%d)=%d\n", x, f(x));
    return 0;
}
```

Diagram illustrating the mapping of C code to memory sections:

- `buf` (array) → `.bss`
- `command` (string) → `.data`
- `n_lines` (variable) → `.bss`
- `n_tries` (variable) → `.data`
- `total` (variable) → `.data`
- `result` (local variable) → `.stack`
- `number_calls` (static variable) → `.bss`
- `x` (local variable) → `.stack`

Text: where the code of the program goes

Data: where global and static variables are:

.data section: global and static initialized variables to non-zero

.bss section: global and static non-initialized variables (or initialized to 0)

Heap: where all dynamically allocated memory is set (e.g. as a result of `malloc()`)

Stack: where all automatic variables exist. Variables that are automatically created and destroyed in methods

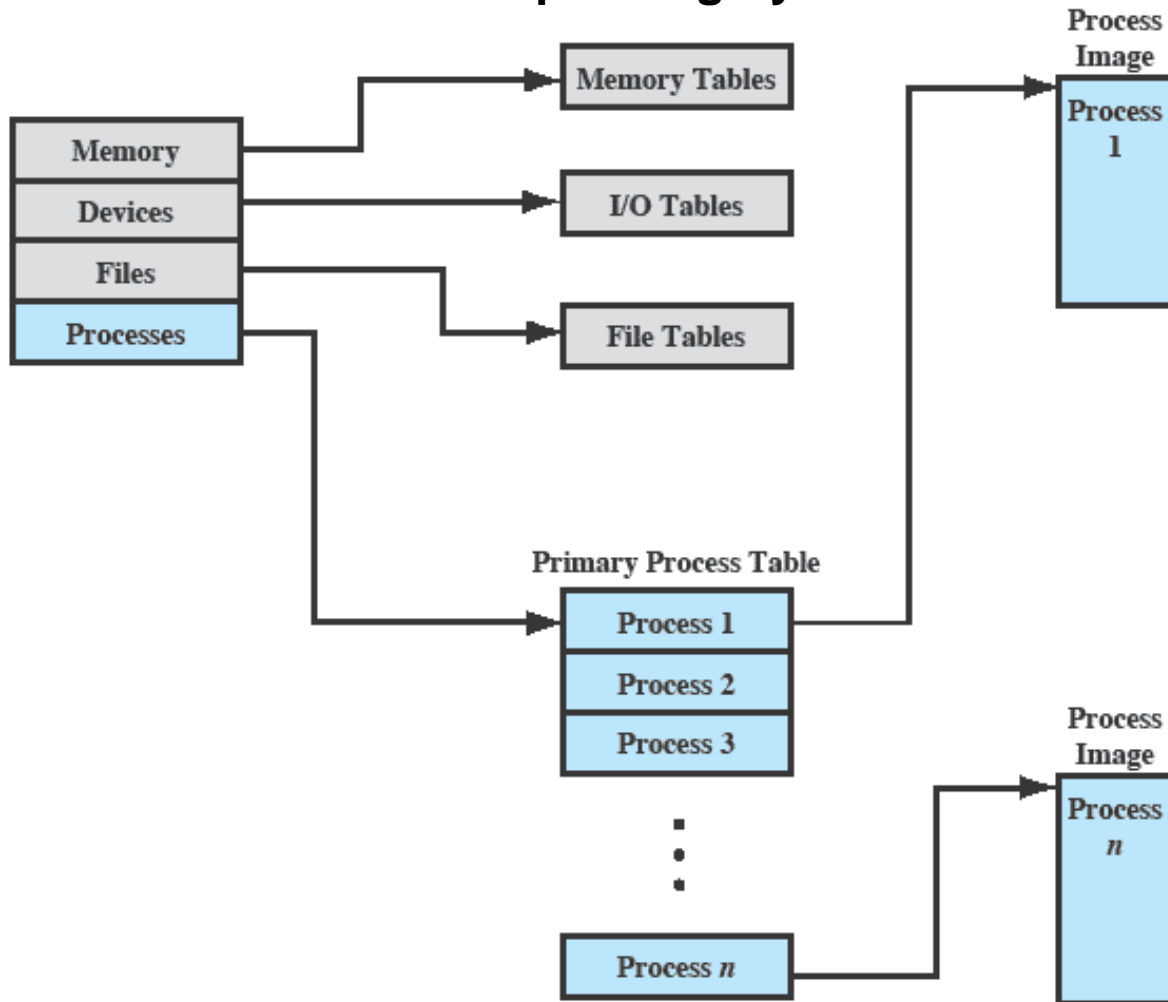
What about the globals:

```
char buf[10*MB] = {0};
```

```
char buf[10*MB] = {1};
```

- **Process**


General Structure of Operating System Control Tables




- **Process Concept**
 - process stack; data section; heap; text
- **Process State**
 - New; Running; Waiting; Ready; Terminated
- **Process Control Block (PCB)**
- **Scheduling Queues**
 - Job Queue
 - Ready Queue
 - Device Queue
- **Schedulers**
 - Long-term scheduler
 - Short-term scheduler
- **Operations in Processes**
 - Process Creation
 - fork(), exec(), wait()
 - Process Termination
 - exit(), wait()
- **References**

- **Process State**

- The state of a process is defined in part by the current activity of that process. Each process may be in one of the following states:
 - **New**: the process is being created.
 - **Running**: instructions are being executed.
 - **Waiting**: the process is waiting for some event to occur (such as an I/O completion or reception of a signal).
 - **Ready**: the process is waiting to be assigned to a processor.
 - **Terminated**: the process has finished execution.



• Process State



William Stallings, “Operating Systems: Internals and Design Principles”

• Process State

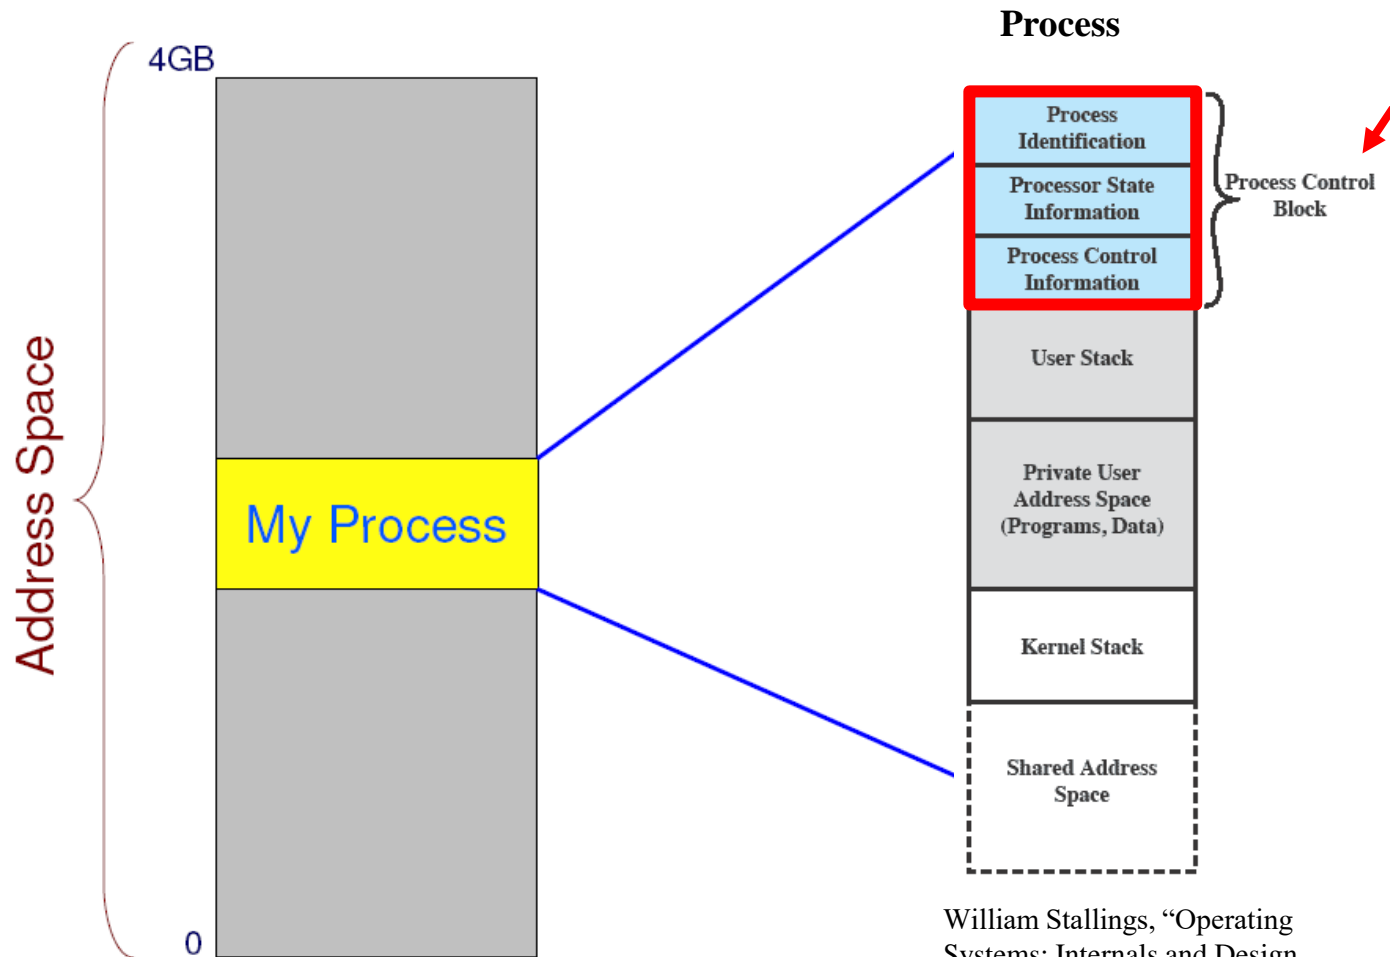
To show the state of a process, use the **ps** command.

```
$ ps -el
F S      UID      PID  PPID   C  PRI   NI  ADDR  SZ  WCHAN  TTY          TIME CMD
0 S      0         1      0    0   75    0    -    155 schedu ?        00:00:04 init
0 S      0         2      1    0   75    0    -      0 contex ?        00:00:00 keventd
0 S      0         3      1    0   94   19    -      0 ksofti ?        00:00:00 ksoftirqd_CPU0
0 S      0         4      1    0   85    0    -      0 kswapd ?        00:00:00 kswapd
0 S      0         5      1    0   85    0    -      0 bdflush ?       00:00:00 bdflush
0 S      0         6      1    0   75    0    -      0 schedu ?        00:00:00 kupdated
0 S      0         7      1    0   85    0    -      0 kinode ?        00:00:00 kinoded
0 S      0         8      1    0   85    0    -      0 md_thr ?       00:00:00 mdrecoveryd
0 S      0        11      1    0   75    0    -      0 schedu ?        00:00:00 kreiserfsd
0 S      0       386      1    0   60  -20    -      0 down_i ?       00:00:00 lvm-mpd
0 S      0       899      1    0   75    0    -    390 schedu ?        00:00:00 syslogd
```

symbol	meaning
S	sleeping
R	running
D	waiting (usually for IO)
T	stopped (suspended) or traced
Z	zombie (defunct)

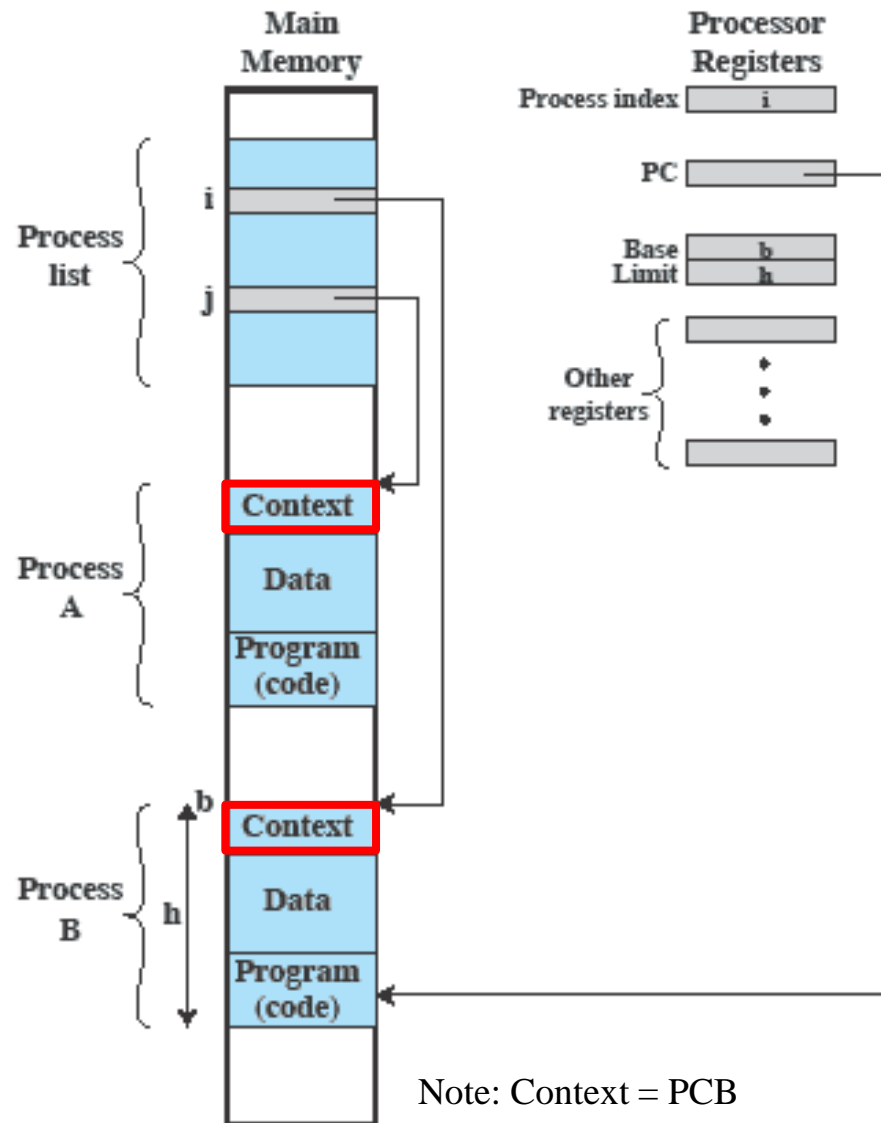
- **Process Concept**
 - process stack; data section; heap; text
- **Process State**
 - New; Running; Waiting; Ready; Terminated
- **Process Control Block (PCB)**
- **Scheduling Queues**
 - Job Queue
 - Ready Queue
 - Device Queue
- **Schedulers**
 - Long-term scheduler
 - Short-term scheduler
- **Operations in Processes**
 - Process Creation
 - fork(), exec(), wait()
 - Process Termination
 - exit(), wait()
- **References**

- **Process Control Block (PBC)**



William Stallings, "Operating Systems: Internals and Design Principles"

- **Process Control Block (PCB)**



- **Process Control Block (PBC)**

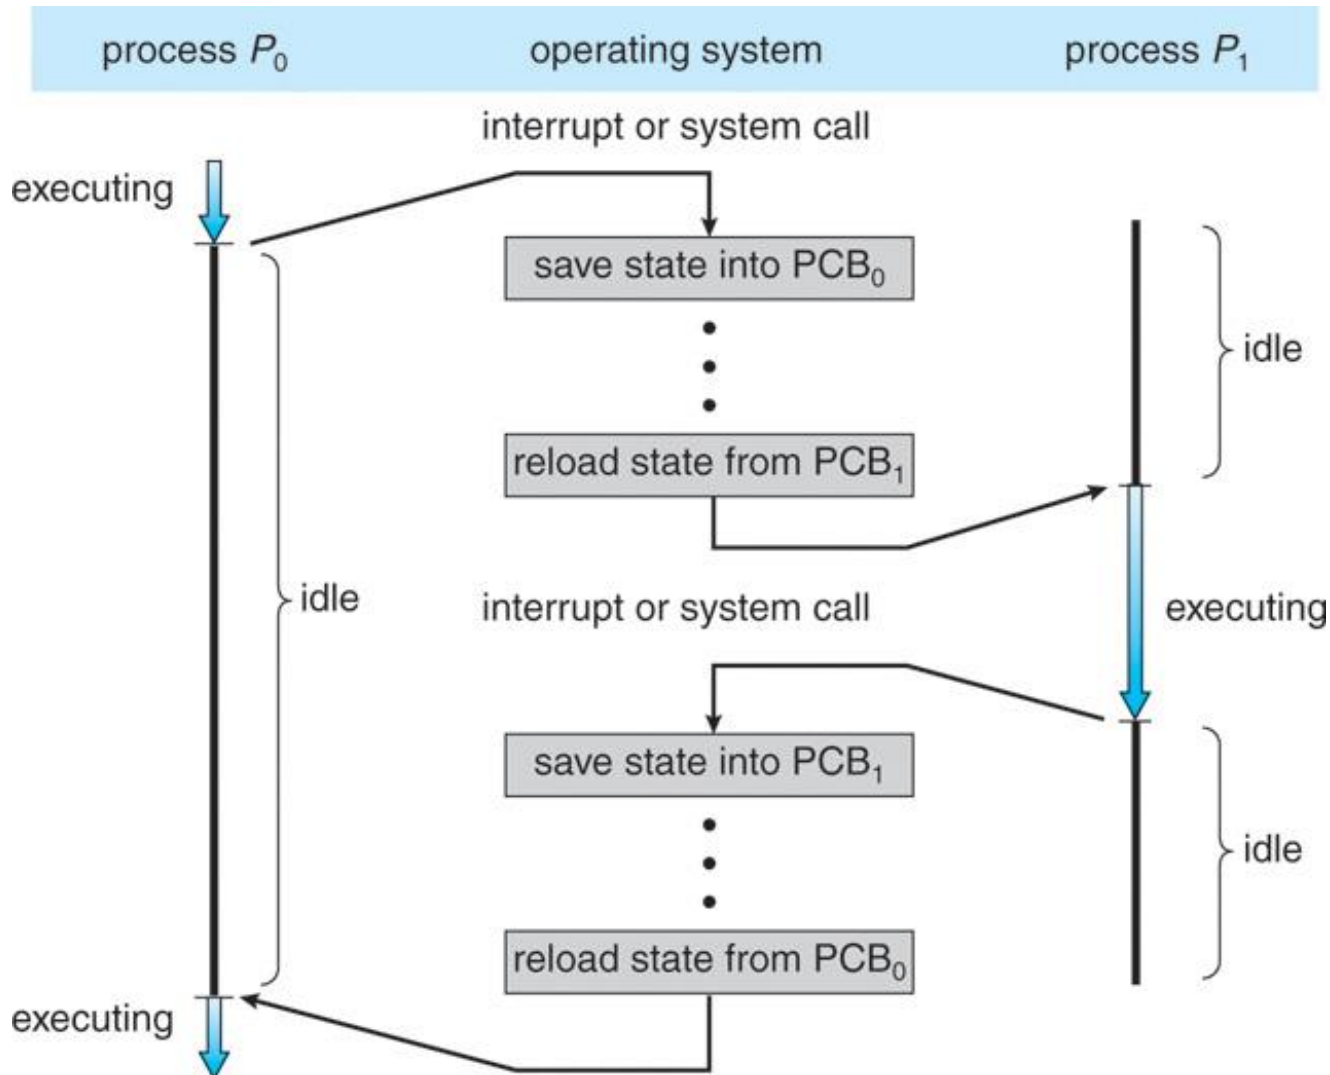
- Each process is represented in the operating system by a *Process Control Block* (PCB) - also called a task control block.

- **Process state:** the state may be new, ready, running, waiting
- **CPU-scheduling information:** includes a process priority, pointers to scheduling queues, and any other scheduling parameters.
- **Program counter:** The counter indicates the address of the next instruction to be executed for this process.
- **Memory-management information:** includes the value of the base and limit registers, the page tables, or the segment tables,
- **CPU registers:** Along with the program counter, this state information must be saved when an interrupt occurs, to allow the process to be continued correctly afterwards
- **I/O status information:** includes the list of I/O devices allocated to the process, a list of open files, and so on.
- **Accounting information:** This information includes the amount of CPU and real time used, time limits, account numbers, job or process numbers, and so on.

PCB



Identifier
State
Priority
Program counter
Memory pointers
Context data
I/O status information
Accounting information
⋮

- Process Control Block (PBC)**



• Example

Dispatcher is a small program which switches the processor from one process to another



1	5000
2	5001
3	5002
4	5003
5	5004
6	5005
Timeout	
7	100
8	101
9	102
10	103
11	104
12	105
13	8000
14	8001
15	8002
16	8003
IO Request	
17	100
18	101
19	102
20	103
21	104
22	105
23	12000
24	12001
25	12002
26	12003

27	12004
28	12005
Timeout	
29	100
30	101
31	102
32	103
33	104
34	105
35	5006
36	5007
37	5008
38	5009
39	5010
40	5011
Timeout	
41	100
42	101
43	102
44	103
45	104
46	105
47	12006
48	12007
49	12008
50	12009
51	12010
52	12011
Timeout	

100 = Starting address of dispatcher program

Shaded areas indicate execution of dispatcher process;
first and third columns count instruction cycles;
second and fourth columns show address of instruction being executed

• Example

5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

(a) Trace of Process A

(b) Trace of Process B

(c) Trace of Process C

5000 = Starting address of program of Process A
8000 = Starting address of program of Process B
12000 = Starting address of program of Process C

1	5000		
2	5001		
3	5002		
4	5003		
5	5004		
6	5005		
----- Timeout			
7	100		
8	101		
9	102		
10	103		
11	104		
12	105		
13	8000		
14	8001		
15	8002		
16	8003		
----- I/O Request			
17	100		
18	101		
19	102		
20	103		
21	104		
22	105		
23	12000		
24	12001		
25	12002		
26	12003		
27	12004		
28	12005		
----- Timeout			
29	100		
30	101		
31	102		
32	103		
33	104		
34	105		
35	5006		
36	5007		
37	5008		
38	5009		
39	5010		
40	5011		
----- Timeout			
41	100		
42	101		
43	102		
44	103		
45	104		
46	105		
47	12006		
48	12007		
49	12008		
50	12009		
51	12010		
52	12011		
----- Timeout			

100 = Starting address of dispatcher program

Shaded areas indicate execution of dispatcher process;
first and third columns count instruction cycles;
second and fourth columns show address of instruction being executed

- **Process Concept**
 - process stack; data section; heap; text
- **Process State**
 - New; Running; Waiting; Ready; Terminated
- **Process Control Block (PCB)**
- **Scheduling Queues**
 - Job Queue
 - Ready Queue
 - Device Queue
- **Schedulers**
 - Long-term scheduler
 - Short-term scheduler
- **Operations in Processes**
 - Process Creation
 - fork(), exec(), wait()
 - Process Termination
 - exit(), wait()
- **References**

• Scheduling Queues

- Job Queue


- As processes enter the system, they are put into a **job queue**, which consists of all processes in the system.

- Ready Queue

- The processes that are residing in main memory and are ready and waiting to execute are kept on a list called the **ready queue**. This queue is generally stored as a linked list.
- A ready-queue header contains pointers to the first and final PCBs in the list. Each PCB includes a pointer field that points to the next PCB in the ready queue.

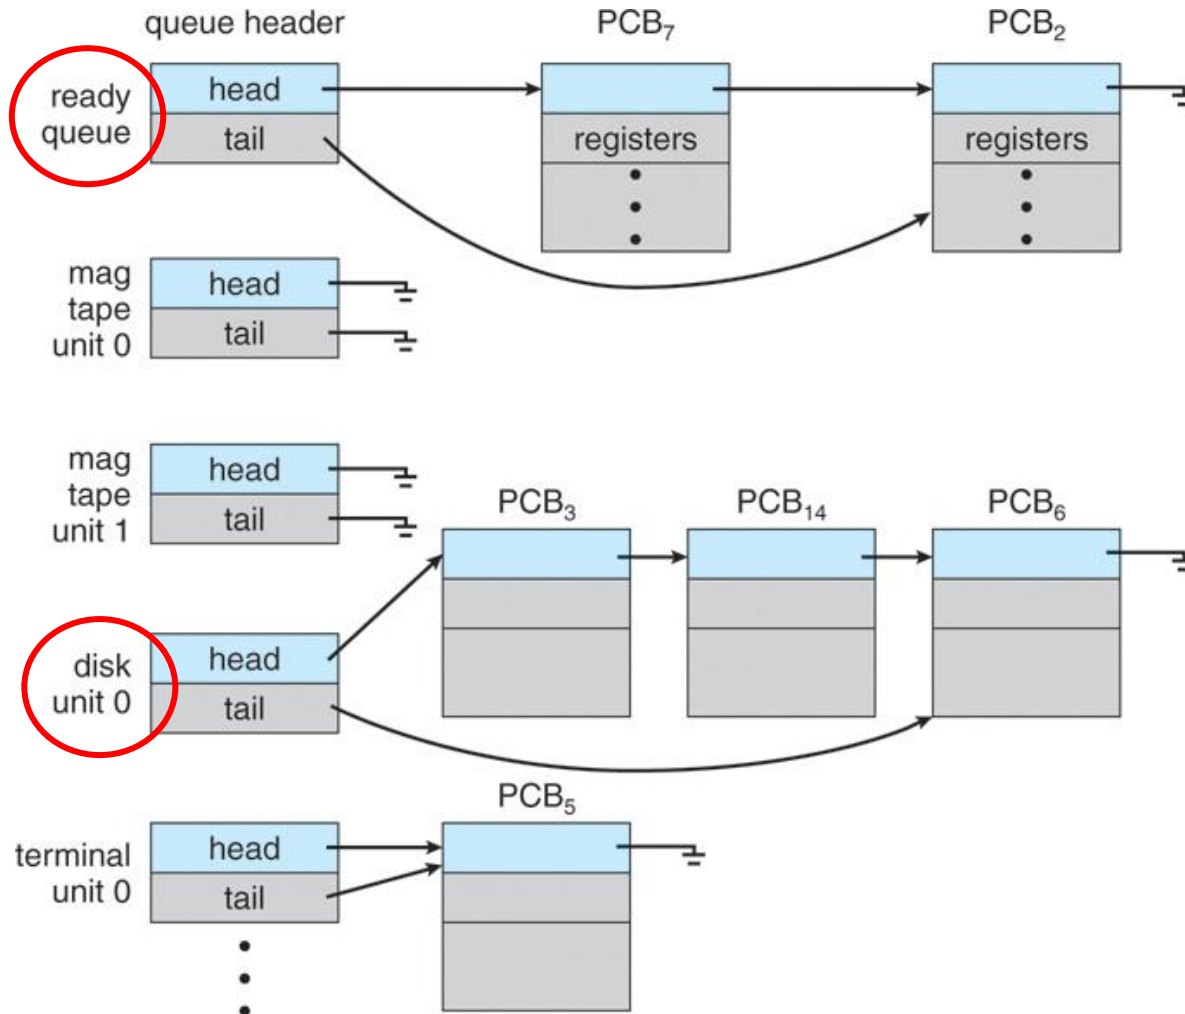
- Device Queue

- The list of processes waiting for a particular I/O device is called a **device queue**. Each device has its own device queue.



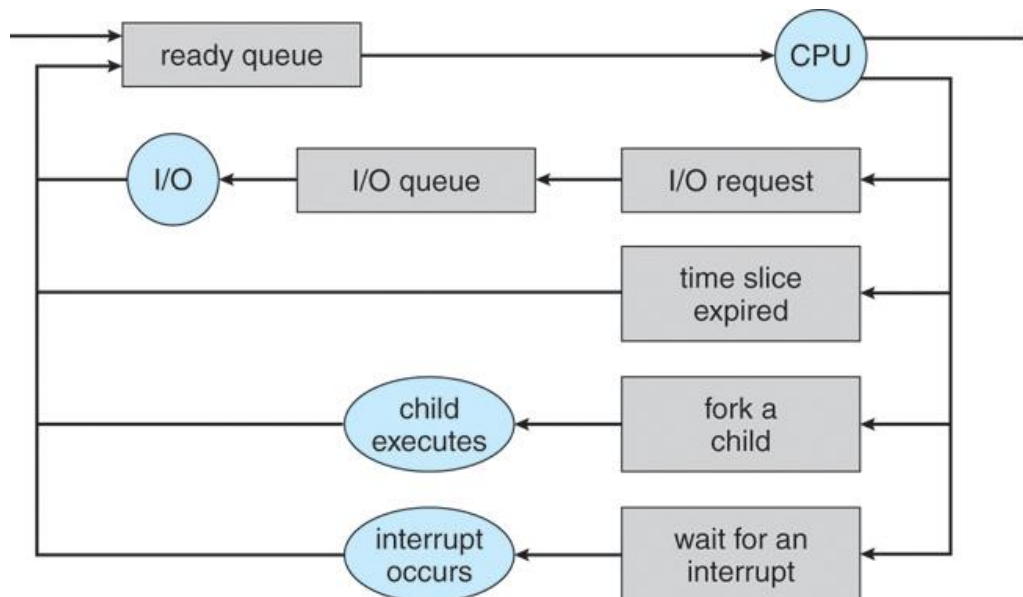
• Scheduling Queues

The ready queue and various I/O device queues.

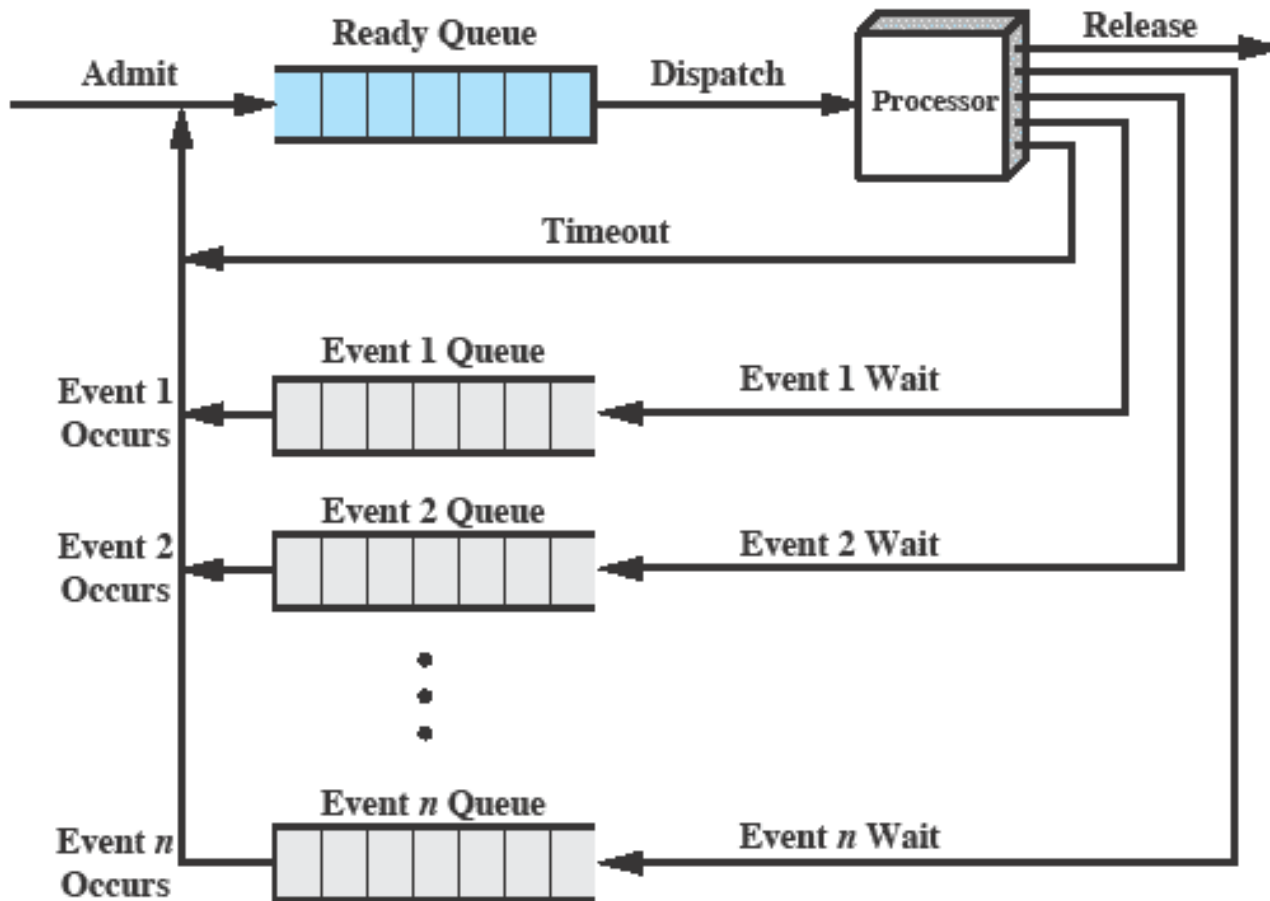


• Queueing diagram:

- A new process is initially put in the ready queue. It waits there until it is selected for execution, or is dispatched.
- Once the process is allocated the CPU and is executing, one of several events could occur:
 - could issue an I/O request and then be placed in an I/O queue.
 - could create a new subprocess and wait for the subprocess's termination.
 - could be removed forcibly from the CPU, as a result of an interrupt, and be put back in the ready queue.



- Queueing diagram:

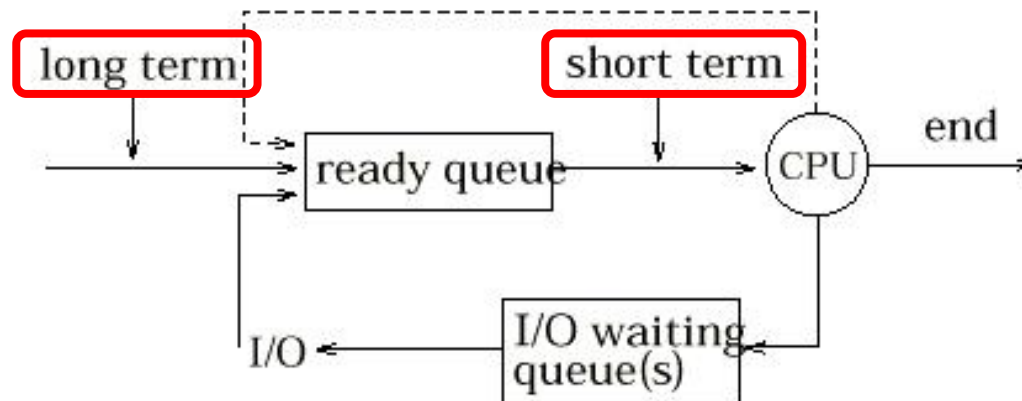


William Stallings, "Operating Systems: Internals and Design Principles"

- **Process Concept**
 - process stack; data section; heap; text
- **Process State**
 - New; Running; Waiting; Ready; Terminated
- **Process Control Block (PCB)**
- **Scheduling Queues**
 - Job Queue
 - Ready Queue
 - Device Queue
- **Schedulers**
 - Long-term scheduler
 - Short-term scheduler
- **Operations in Processes**
 - Process Creation
 - fork(), exec(), wait()
 - Process Termination
 - exit(), wait()
- **References**

- **Schedulers**

- A process migrates among the various scheduling queues throughout its lifetime. The selection from these queues is carried out by the appropriate scheduler.
- **long-term scheduler:**
 - processes are spooled to a mass-storage device (typically a disk), where they are kept for later execution. The long-term scheduler will select processes from this pool and loads them into memory for execution.
- **short-term scheduler**, or CPU scheduler,
 - selects from among the processes that are ready to execute and allocates the CPU to one of them.
- It is important that the long-term scheduler select a good process mix of I/O-bound and CPU-bound processes

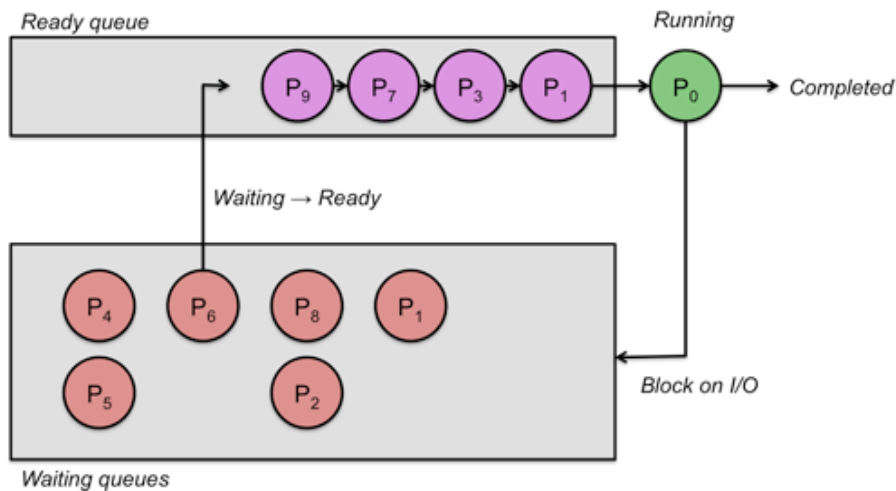


• Schedulers

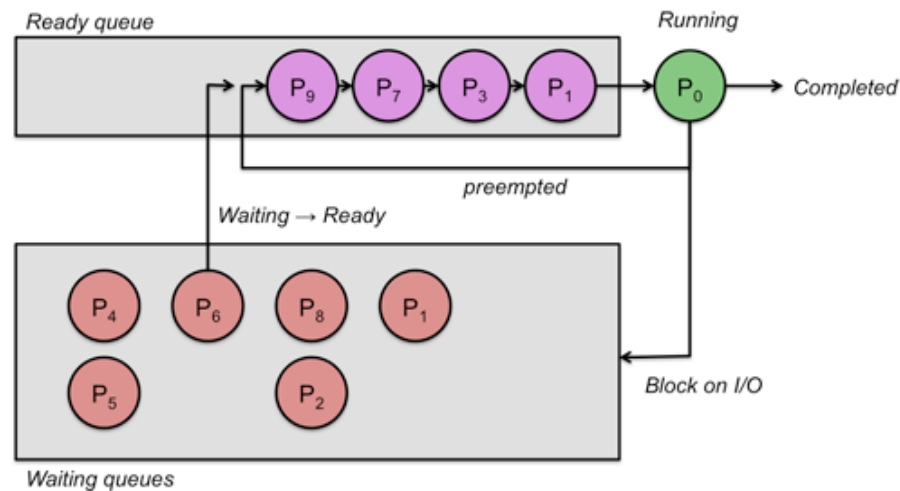
- What's the best scheduler? It depends:

- Server Operating Systems
- User Operating Systems
- Real-time Operating Systems
- Embedded Operating Systems
- Multimedia Operating Systems
-

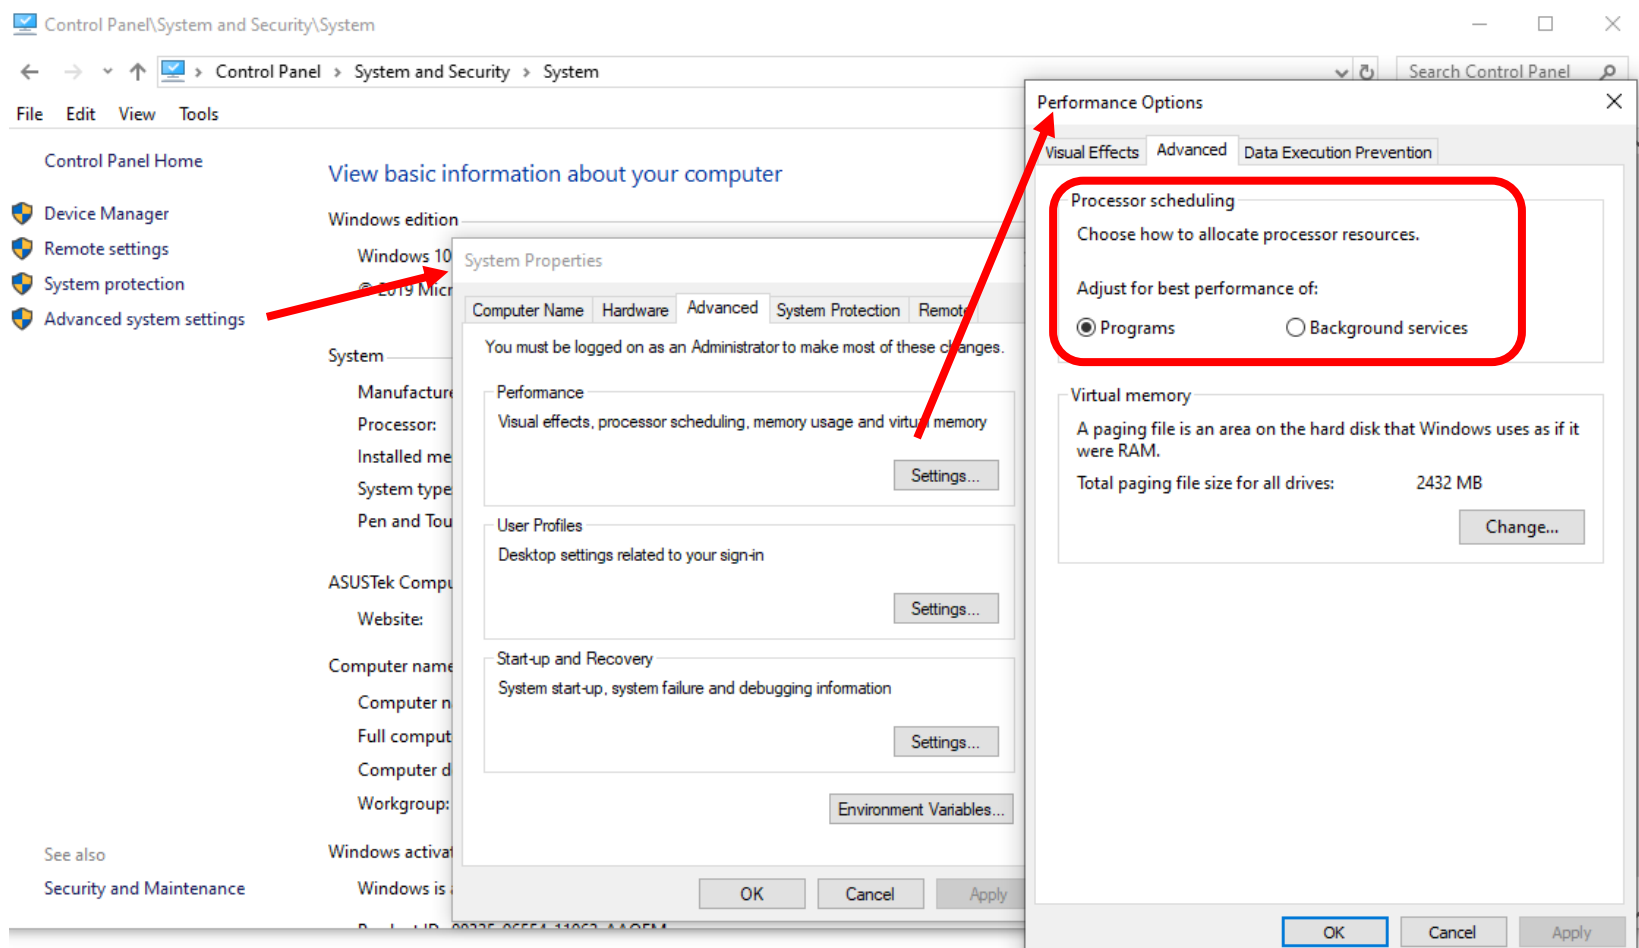
First Come First Served



Round Robin



- **Schedulers**
 - **Even in Windows!**






Control Panel » System Security » System » Advanced System
Settings » Advanced » Performance » Settings » Advanced

- **Schedulers**

- **Context switch**

- when changing of running process, the system needs to save the current context of the process currently running on the CPU so that it can restore that context when resuming it
- The **context** is represented in the PCB of the process, and it includes the value of the CPU registers, the process state, and memory-management information.



- **Schedulers**
- **Context switch**
 - A process switch also occurs when the OS gains control from the currently running process. Possible events giving OS control are:

Mechanism	Cause	Use
Interrupt	External to the execution of the current instruction	Reaction to an asynchronous external event
Trap	Associated with the execution of the current instruction	Handling of an error or an exception condition
Supervisor call	Explicit request	Call to an operating system function

William Stallings, “Operating Systems: Internals and Design Principles”

- **Schedulers**
- **Why Event Driven... “A Human perspective”**


Characteristic			Scaled to Human Time	
<i>CPU frequency</i>	<i>2GHz</i>			
Processor Cycle Time	0.5 ns		1 s	
L2 cache access	10 ns		20 s	
Memory access	80 ns		160 s (2.6 mins)	
Thread context switch	5000 ns (5us)		10000 s (2.7 hours)	
Disk access	8000000 ns (8ms)		16000000 s (185 days)	
Process quantum	100000000 ns (100ms)		200000000 s (6.3 years)	
In blue ► Things improving very fast				
In orange ► Things improving to a degree				
In red ► Things not really improving				

- **Process Concept**
 - process stack; data section; heap; text
- **Process State**
 - New; Running; Waiting; Ready; Terminated
- **Process Control Block (PCB)**
- **Scheduling Queues**
 - Job Queue
 - Ready Queue
 - Device Queue
- **Schedulers**
 - Long-term scheduler
 - Short-term scheduler
- **Operations in Processes**
 - Process Creation
 - fork(), exec(), wait()
 - Process Termination
 - exit(), wait()
- **References**

• Process Creation

- A process may create several new processes, during the course of execution.
- The creating process is called a **parent process**, and the new processes are called the **children** of that process.
- Each of these new processes may in turn create other processes, forming a **tree** of processes.

Process Hierarchies



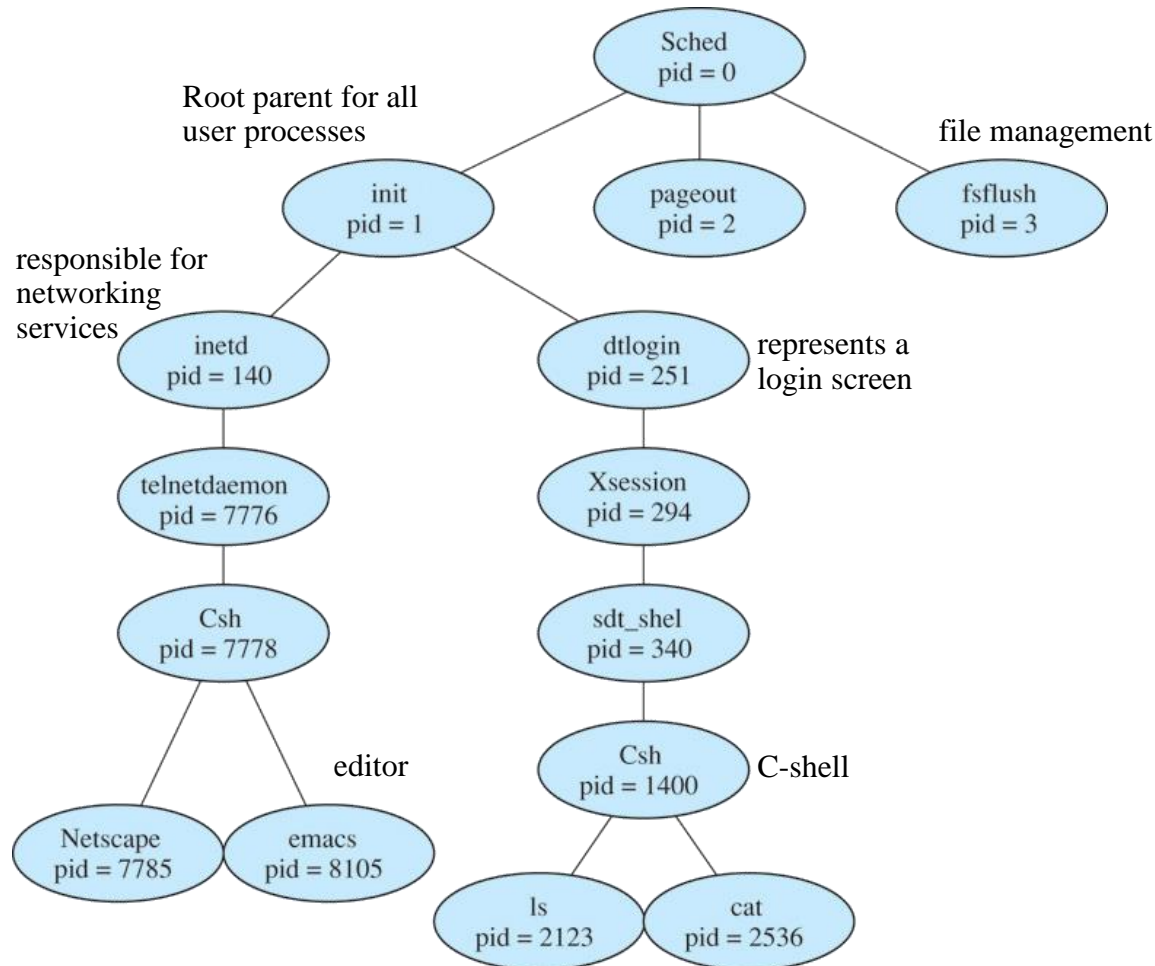
In Unix the **fork** system call creates a child process, and the **exit** system call terminates the current process.

1) After a `fork` both parent and child keep running (indeed they have the same program text) and each can fork off other processes.

2) A process can choose to **wait** for children to terminate. For example, if C issued a `wait()` system call it would block until G finished.

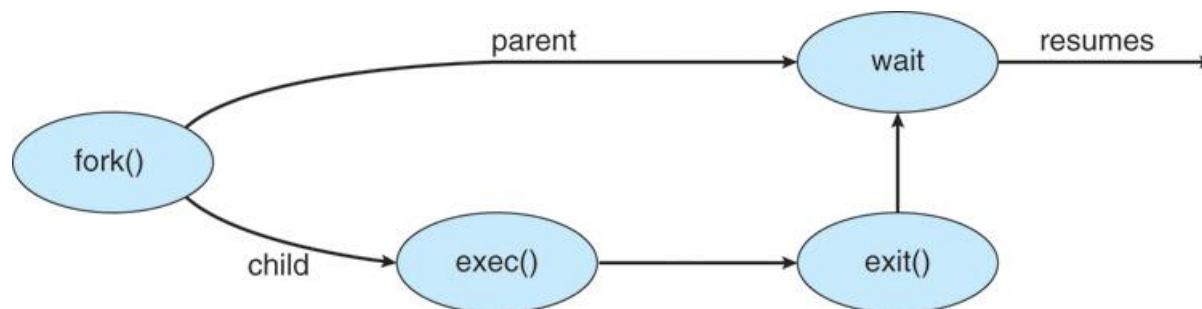
• Process Creation

A tree of processes on a typical Solaris system.



• Process Creation

- Most operating systems (including UNIX and the Windows family of operating systems) identify processes according to a unique **Process Identifier** (or **PID**), which is typically an integer number.
- When a process creates a new process, **two possibilities exist in terms of execution**:
 - 1. The parent continues to execute concurrently with its children.
 - 2. The parent waits until some or all of its children have terminated.



Abraham Silberschatz, "Operating System Concepts"

• Process Creation

```
#include <sys/types.h>
#include <stdio.h>
```

```
int main(){
    pid_t pid;
    // fork a child process
    pid = fork();
    if (pid < 0) { // error occurred
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { //child process
        execlp("/bin/ls", "ls", NULL);
    }
    else { //parent process
        //parent will wait for the child to complete
        wait(NULL);
        printf("Child Complete");
        exit(0);
    }
}
```

Fork: this system call creates a new process.

The return code for the fork is:

- **zero** for the new (child) process;
- **nonzero process identifier** of the child for the parent process.

- **Process Creation**

- **Fork:** this system call creates a new process which consists of a copy of the address space of the original process. Both processes (the parent and the child) continue execution at the instruction after the fork, with a difference. The return code for the fork is:
 - **zero** for the new (child) process;
 - **nonzero process identifier** of the child for the parent process.
- **Exec:** this system call can be used to replace the process's memory space with a new program:
 - loads a binary file into memory (destroying the memory image of the program containing the exec system call) and starts its execution.
- **Wait:** if it has nothing else to do while the child runs, the parent process can issue a wait system call to move itself off the ready queue until the termination of the child.

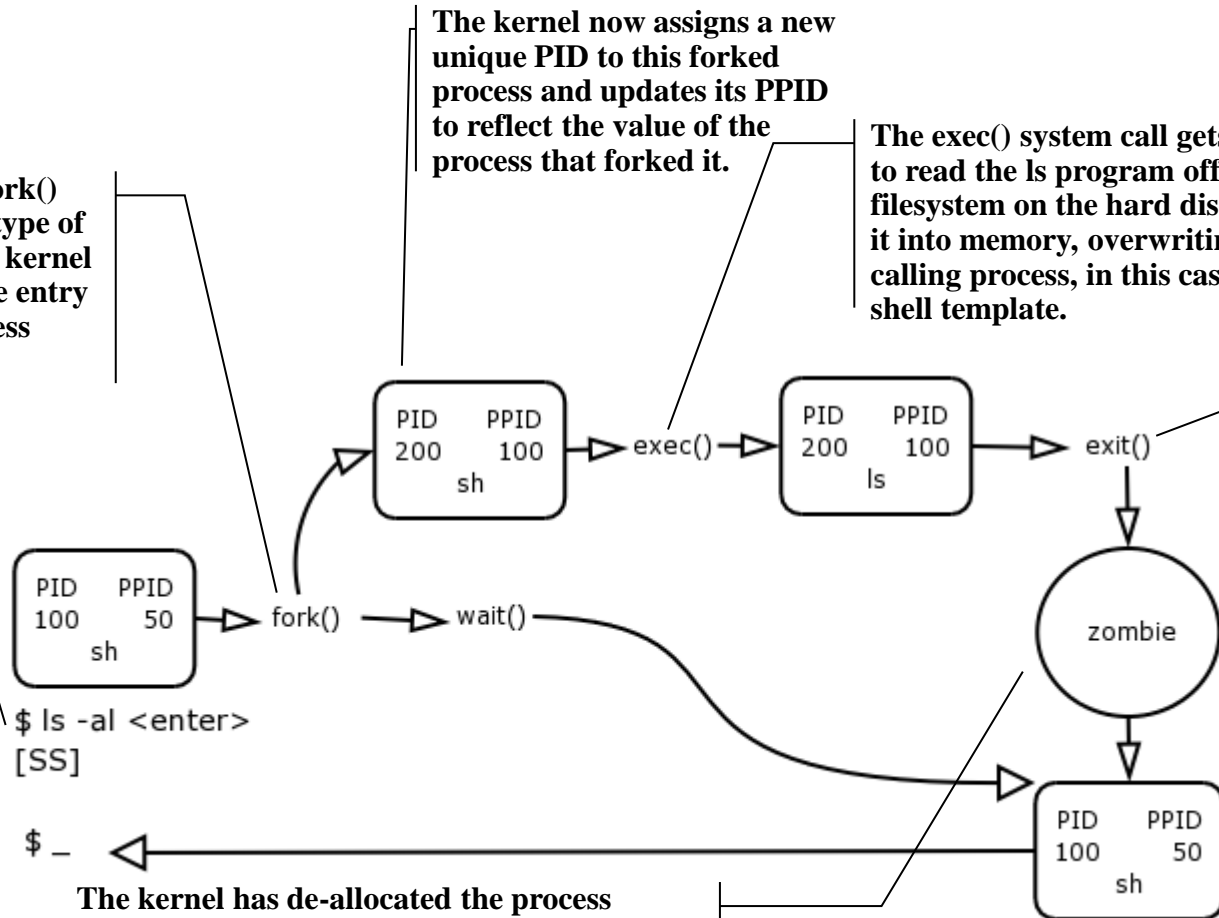
• Process Creation

What it does first is to issue a `fork()` system call. Which performs a type of cloning operation. This gets the kernel to copy an existing process table entry to a next empty slot in the process table. This effectively creates a template for a new process.

The kernel now assigns a new unique PID to this forked process and updates its PPID to reflect the value of the process that forked it.

The `exec()` system call gets the kernel to read the `ls` program off the filesystem on the hard disk and place it into memory, overwriting the calling process, in this case the child shell template.

The `exit()` of the child actually causes the return of the `wait()` system call, which ends the pausing of the parent process, so that it can now continue running.



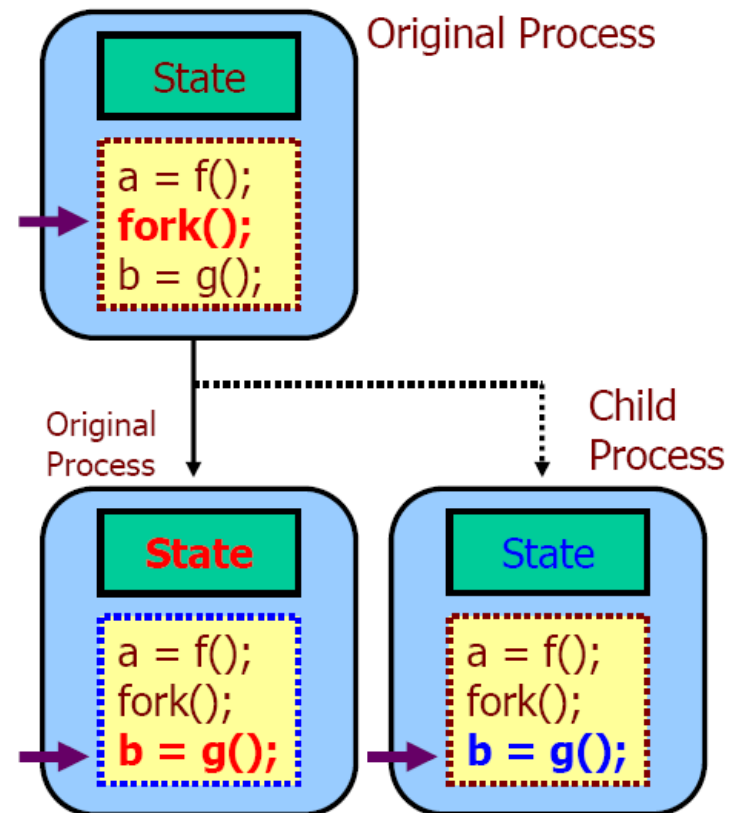
The kernel has de-allocated the process memory, however its process table entry still exists. It is the job of the parent to inform the kernel that it has finished working with the child, and that the kernel can now remove the process table entry from the child

• Process Creation

- Process Model Unix

- Process creation in Unix is based on spawning child processes which inherit all the characteristics of their fathers:

- Variables, program counter, open files, etc.
 - Spawning a process is done using the **fork()** system call
- After forking, each process will be executing having different variables and different state.



• Process Creation

"init" is the first process run by the kernel at bootup: this behavior is hard-coded in the kernel. It is "init's" job to start up various child processes to get the system to a usable state

```
$ ps -ef | less
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	0	1	0	0	75	0	-	155	schedu	?	00:00:04	init
0	S	0	2	1	0	75	0	-	0	contex	?	00:00:00	keventd
0	S	0	3	1	0	94	19	-	0	ksofti	?	00:00:00	ksoftirqd_CPU0
0	S	0	4	1	0	85	0	-	0	kswapd	?	00:00:00	kswapd
0	S	0	5	1	0	85	0	-	0	bdfus	?	00:00:00	bdfus
0	S	0	6	1	0	75	0	-	0	schedu	?	00:00:00	kupdated
0	S	0	7	1	0	85	0	-	0	kinode	?	00:00:00	kinoded
0	S	0	8	1	0	85	0	-	0	md_thr	?	00:00:00	mdrecoveryd
0	S	0	11	1	0	75	0	-	0	schedu	?	00:00:00	kreiserfsd
0	S	0	386	1	0	60	-20	-	0	down_i	?	00:00:00	lvm-mpd
0	S	0	899	1	0	75	0	-	390	schedu	?	00:00:00	syslogd
0	S	0	902	1	0	75	0	-	593	syslog	?	00:00:00	klogd

the **PPID** column shows the parent process of the program.

• Process Creation

Process Explorer - Sysinternals: www.sysinternals.com [LAPTOP-ID2IHJMH\LuisVeloso]

File Options View Process Find Users Help

Process	CPU	Private Bytes	CPU	Description	Company Name
Registry		8,060 K	4.31%		
System Idle Process	95.28	60 K	1.12%	procexp64.exe:17368	
System	0.31	192 K	19:01:42		
Interrupts	0.35	0 K	0 K	n/a Hardware Interrupts and DPCs	
smss.exe		1,152 K	476 K	460	
Memory Compression		228 K	10,500 K	2972	
csrss.exe		1,844 K	2,324 K	620	
wininit.exe		1,396 K	1,684 K	724	
services.exe		6,032 K	8,764 K	800	
svchost.exe		912 K	1,288 K	932 Host Process for Windows S...	Microsoft Corporation
svchost.exe	< 0.01	22,356 K	32,712 K	956 Host Process for Windows S...	Microsoft Corporation
WmiPrvSE.exe		9,704 K	14,268 K	8040	
StartMenuExperience...	0.01	34,752 K	76,304 K	9612	
RuntimeBroker.exe		9,784 K	20,052 K	10000 Runtime Broker	Microsoft Corporation
SearchUI.exe	Susp...	116,548 K	177,304 K	9816 Search and Cortana applicati...	Microsoft Corporation
RuntimeBroker.exe		17,724 K	39,752 K	10488 Runtime Broker	Microsoft Corporation
SettingSyncHost.exe		16,192 K	22,368 K	11060 Host Process for Setting Syn...	Microsoft Corporation
YourPhoneServer.exe		3,204 K	6,268 K	5460	
RuntimeBroker.exe		6,612 K	16,744 K	10784 Runtime Broker	Microsoft Corporation
ApplicationFrameHost...		20,204 K	30,112 K	13716 Application Frame Host	Microsoft Corporation
SystemSettings.exe	Susp...	22,304 K	784 K	5300 Settings	Microsoft Corporation
MicrosoftEdge.exe	Susp...	24,392 K	2,456 K	10016 Microsoft Edge	Microsoft Corporation
browser_broker.exe		1,696 K	2,424 K	4464 Browser_Broker	Microsoft Corporation
RuntimeBroker.exe		1,608 K	2,376 K	6240 Runtime Broker	Microsoft Corporation
MicrosoftEdgeSH...	Susp...	5,736 K	1,512 K	7364 Microsoft Edge Web Platform	Microsoft Corporation
MicrosoftEdgeCP.exe	Susp...	7,584 K	5,580 K	8200 Microsoft Edge Content Proc...	Microsoft Corporation
LockApp.exe	Susp...	13,772 K	34,028 K	1688 LockApp.exe	Microsoft Corporation
RuntimeBroker.exe		8,584 K	24,300 K	10436 Runtime Broker	Microsoft Corporation
dllhost.exe		4,060 K	7,784 K	12440 COM Surrogate	Microsoft Corporation

CPU Usage: 4.72% Commit Charge: 40.27% Processes: 222 Physical Usage: 36.96%

Process Explorer from SysInternals for Windows

<https://docs.microsoft.com/pt-pt/sysinternals/>

- **Process Termination**

- A process terminates when it finishes executing its final statement and asks the operating system to delete it by using the ***exit()*** system call.
 - the process may return a status value (typically an integer) to its parent process via the ***wait()*** system call.
 - all the resources of the process (including physical and virtual memory, open files, and I/O buffers) are deallocated by the operating system.
- A process can cause the termination of another process via an appropriate system call:
 - usually, such a system call can be invoked only by the parent of the process that is to be terminated. Otherwise, users could arbitrarily kill each other's jobs.

- System Calls**

Process Management		
Posix	Win32	Description
fork	CreateProcess	<u>Clone</u> current process
exec		<u>Replace</u> current process
wait	WaitForSingleObject	Wait for a child to terminate.
exit	ExitProcess	Terminate current process & return status
File Management		
Posix	Win32	Description
open	CreateFile	Open a file & return descriptor
close	CloseHandle	Close an open file
read	ReadFile	Read from file to buffer
write	WriteFile	Write from buffer to file
lseek	SetFilePointer	Move file pointer
stat	GetFileAttributesEx	Get status info

- Abraham Silberschatz, "Operating System Concepts", 10th Edition, Wiley, 2018
- William Stallings, "Operating Systems: Internals and Design Principles", 9th Edition, Pearson, 2017
- Andrew S. Tanenbaum, "Modern Operating Systems", 4th Edition, Pearson Education, 2014